

Machine Learning: Models and Applications

Lecture 7

Lecturer: Xinchao Wang

Outline

- Linear Regression
 - Regression technique
- Logistic Regression (called regression, but in fact, classification)
 - Classification technique
- Gradient Descent/Ascent
 - Descent -> min.
 - Ascent -> max.
 - Optimization technique
- Perceptron
 - Classification technique

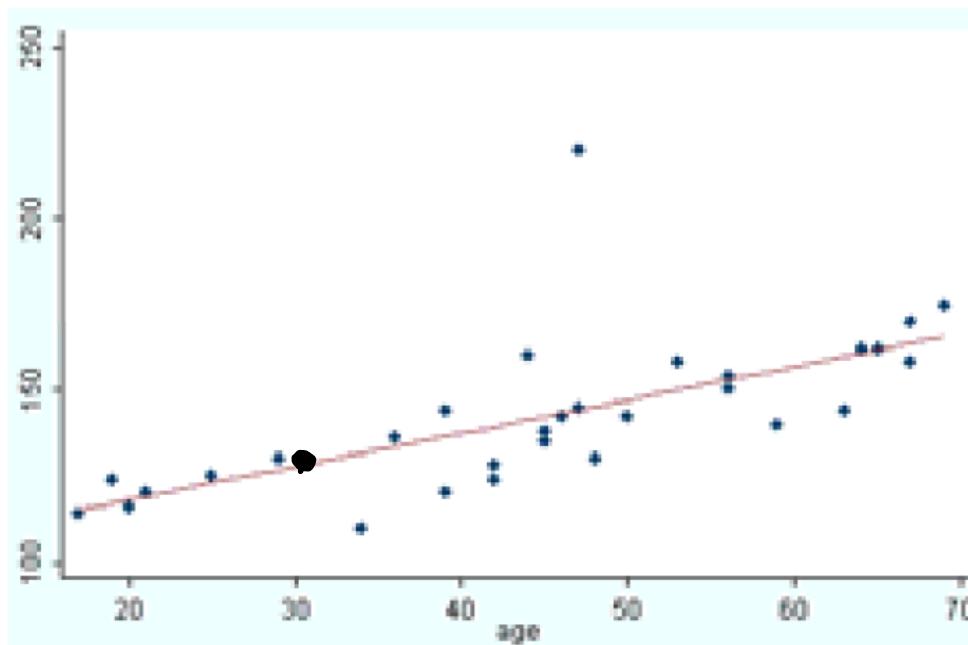
Linear Regression

Simple Linear Regression

- How does a single variable of interest relate to another (single) variable?
 - Y = outcome variable (response, dependent...)
 - X = explanatory variable (predictor, feature, independent...)
- Data: n pairs of continuous observations $(X_1, Y_1), \dots, (X_n, Y_n)$

Example

- How does systolic blood pressure (SBP) relate to age?

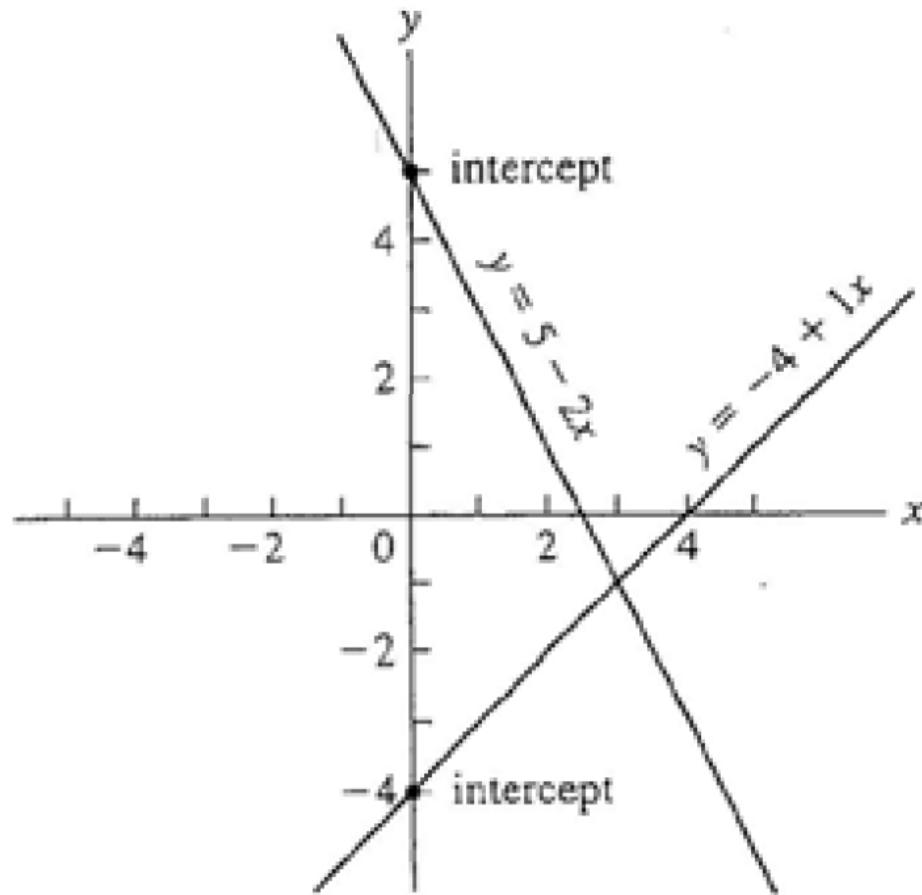


- Graph suggests that Y relates to X in an approximately linear way

Regression: Step by Step

1. Assume a linear model: $Y = \beta_0 + \beta_1 X$
2. Find the line which “best” fits the data, i.e. estimate parameters β_0 and β_1
3. Does variation in X help describe variation in Y ?
4. Check assumptions of model
5. Draw inferences and make predictions

Straight-line Plots



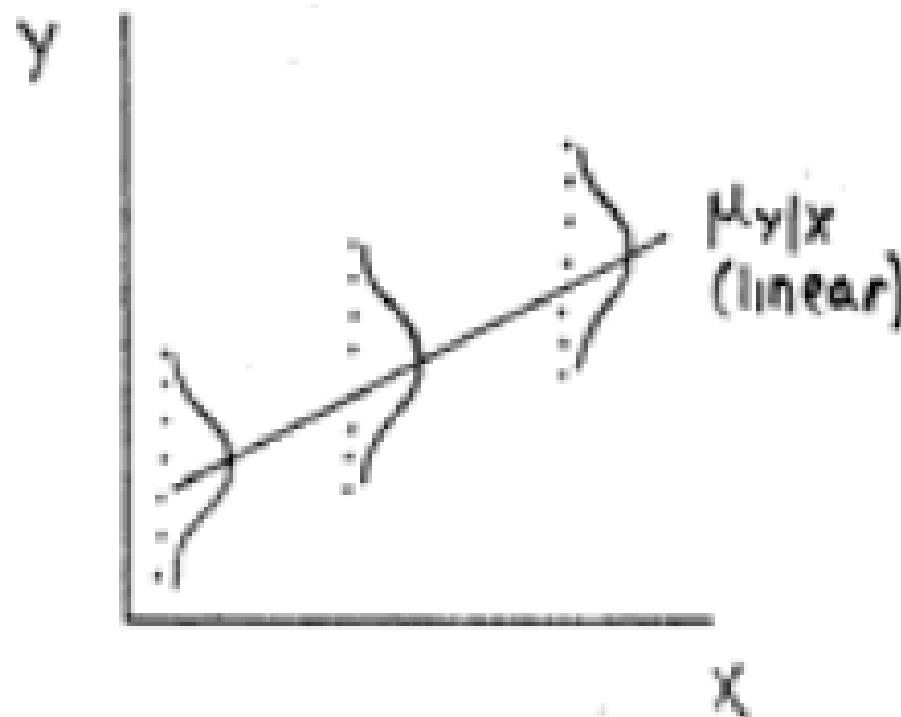
Assumptions of Linear Regression

- Five basic assumptions
1. Existence: for each fixed value of X , Y is a random variable with finite mean and variance
 2. Independence: the set of Y_i are independent random variables given X_i

Assumptions of Linear Regression

3. Linearity: the mean value of Y is a linear function of X

$$\mu_{Y|X} = E[Y|X] = \beta_0 + \beta_1 X$$



Assumptions of Linear Regression

4. Homoscedasticity: the variance of Y is the same for any X
5. Normality: For each fixed value of X, Y has a normal distribution (by assumption 4, σ^2 does not depend on X)

$$Y \sim N(\mu_{Y|X}, \sigma^2)$$

Formulation

$$Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$$

$\varepsilon_i \sim N(0, \sigma^2)$ independent

$\Rightarrow Y_i$ are $N(\beta_0 + \beta_1 X_i, \sigma^2)$ given X_i

$\Rightarrow E(Y_i | X_i) = \beta_0 + \beta_1 X_i$

$Var(Y_i | X_i) = \sigma^2$ = variability of Y_i about $\mu_{Y|X_i}$

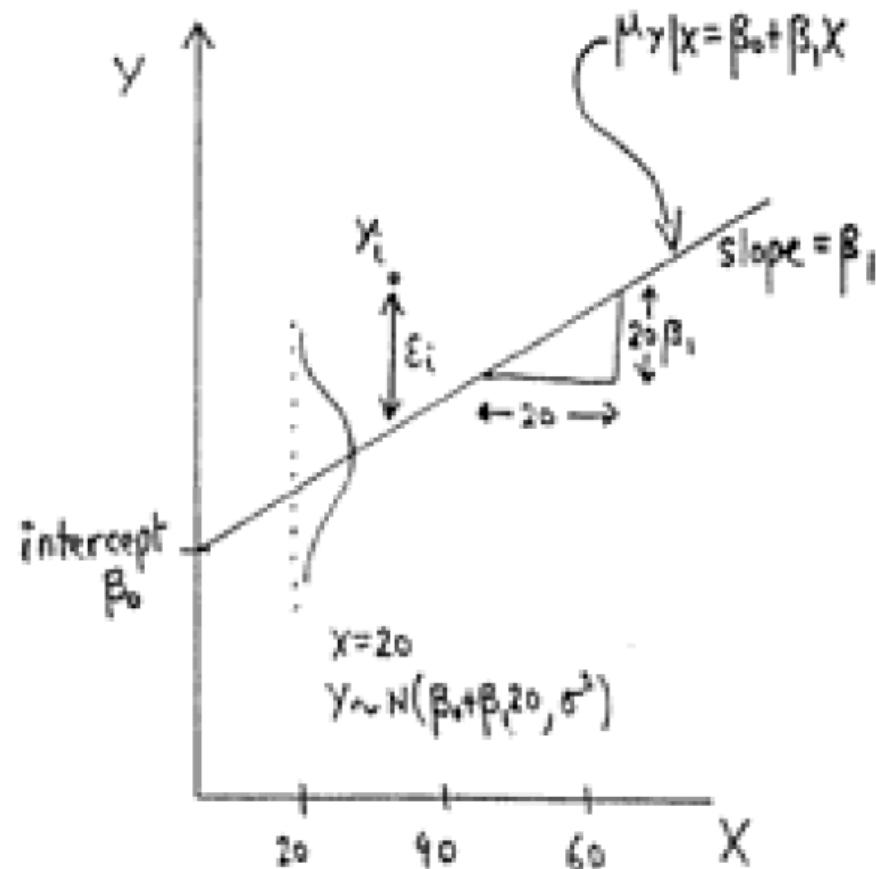
- Y_i are linear function of X_i plus some random error

ε_i = error $(= Y_i - \mu_{Y|X_i})$

$\hat{\varepsilon}_i = Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i$ is called the "residual" for Y_i

$\hat{\varepsilon}_i = Y_i - \hat{Y}_i$

Linear Regression



Estimating β_0 and β_1

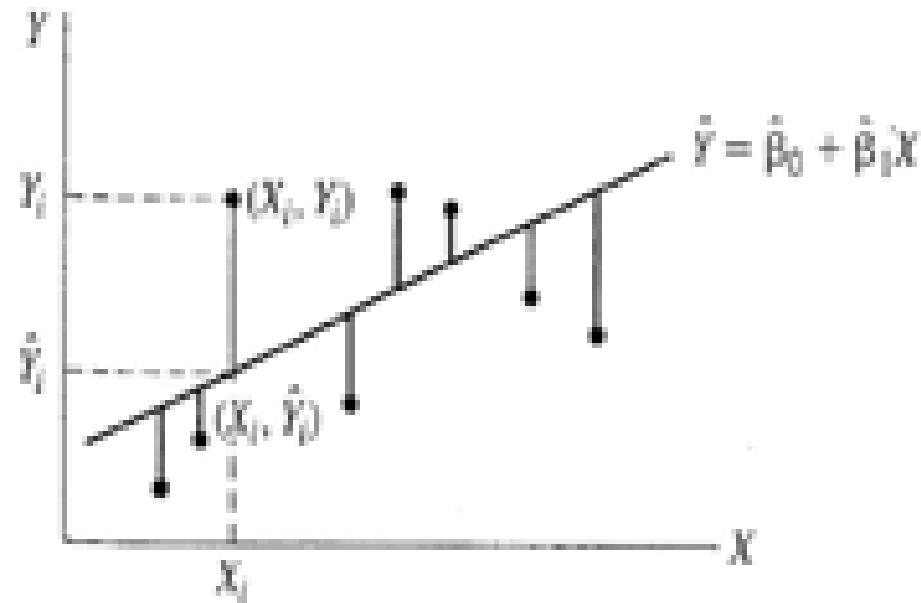
- Find “best” line
- Criterion for “best”: estimate β_0 and β_1 to minimize:

$$\begin{aligned} & \sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)^2 \\ &= \sum_{i=1}^n (Y_i - \hat{Y}_i)^2 = \sum_{i=1}^n \hat{\varepsilon}_i^2 \end{aligned}$$

- This is the residual sum of squares, sum of squares due to error, or sum of squares about regression line
- Least Squares estimator

Rationale for LS Estimates

- $\hat{\varepsilon}^2$ measures the “deviance” of Y_i from the estimated model
- The “best” model is the one from which the data deviate the least



Least Squares Estimators

- Taking derivatives with respect to β , we obtain

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{s_{xy}}{s_x^2}$$

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$$

Example: SBP/age data

$$\bar{X} = 45.13 \quad \bar{Y} = 142.53$$

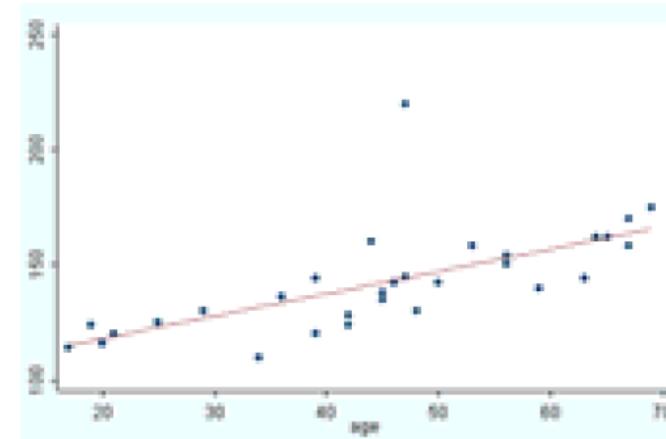
$$S_x^2 = 233.91 \quad S_y^2 = 509.91 \quad S_{xy} = 227.10$$

$$\hat{\beta}_1 = \frac{S_{xy}}{S_x^2} = \frac{227.10}{233.91} = 0.97$$

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X} = 142.53 - 0.97(45.13) = 98.71$$

$$\hat{y} = 98.71 + .97x$$

.97 mm Hg \uparrow for every
1 yr \uparrow in age



Using the Model

- Using the parameter estimates, our best guess for any Y given X is

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X$$

- Hence at \bar{X}

$$\hat{\beta}_0 + \hat{\beta}_1 \bar{X} = (\bar{Y} - \hat{\beta}_1 \bar{X}) + \hat{\beta}_1 \bar{X} = \bar{Y}$$

- Also $\hat{Y} = \bar{Y} + \hat{\beta}_1 (X - \bar{X})$

Regression and Correlation Coefficient

Regression
Coefficient

Correlation
Coefficient

$$\hat{\beta}_1 = \frac{S_{xy}}{S_x^2} \text{ and } r_{xy} = \frac{S_{xy}}{S_x S_y}$$

$$\Rightarrow r_{xy} = \hat{\beta}_1 \cdot \frac{S_x}{S_y}$$

Example

Suppose we have:

x	y
1	2
2	4
-1	0
3	4
4	6

$$n = 5$$

$$\bar{x} = \frac{1+2-1+3+4}{5} = 1.8$$

$$\bar{y} = \frac{2+4+0+4+6}{5} = 3.2$$

Calculate r :

$$r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}}$$

$x_i - \bar{x}$	$y_i - \bar{y}$	$(x_i - \bar{x})(y_i - \bar{y})$
1-1.8 = -0.8	2-3.2 = -1.2	(-0.8)(-1.2) = 0.96
2-1.8 = 0.2	4-3.2 = 0.8	(0.2)(0.8) = 0.16
-1-1.8 = -2.8	0-3.2 = -3.2	(-2.8)(-3.2) = 8.96
3-1.8 = 1.2	4-3.2 = 0.8	(1.2)(0.8) = 0.96
4-1.8 = 2.2	6-3.2 = 2.8	(2.2)(2.8) = 6.16
0	0	17.2

Example

$$\sum (x_i - \bar{x})^2 = (-0.8)^2 + (0.2)^2 + (-2.8)^2 + (1.2)^2 + (2.2)^2 = 14.8$$

$$\sum (y_i - \bar{y})^2 = (-1.2)^2 + (0.8)^2 + (-3.2)^2 + (0.8)^2 + (2.8)^2 = 20.8$$

$$s_x = \sqrt{\frac{14.8}{4}} = 1.92$$

$$s_y = \sqrt{\frac{20.8}{4}} = 2.28$$

$$r = \frac{17.2}{\sqrt{(14.8)(20.8)}} = 0.98$$

$$\hat{\beta}_1 = r \cdot \frac{s_y}{s_x} = (0.98) \left(\frac{2.28}{1.92} \right) = 1.16$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} = 3.2 - (1.16)(1.8) = 1.11$$

Estimated regression line is $\hat{y}_i = 1.11 + 1.16x_i$

Example

Fitted values are the $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$

$$\hat{y}_1 = 1.11 + 1.16(1) = 2.27 \quad \hat{y}_4 = 1.11 + 1.16(3) = 4.59$$

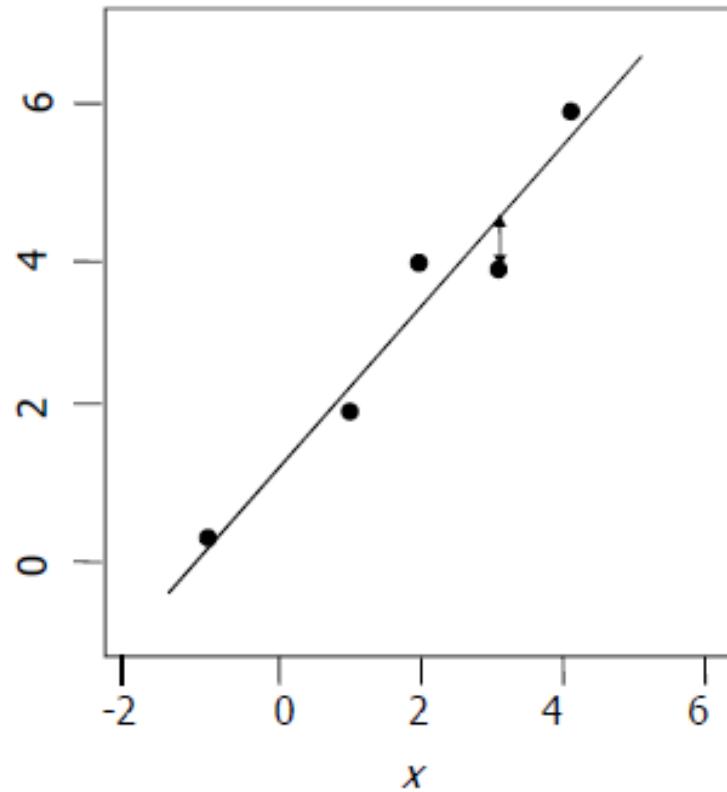
$$\hat{y}_2 = 1.11 + 1.16(2) = 3.43 \quad \hat{y}_5 = 1.11 + 1.16(4) = 5.75$$

$$\hat{y}_3 = 1.11 + 1.16(-1) = -0.05$$

Residuals =

$y_i - \hat{y}_i = \hat{\varepsilon}_i$	$(y_i - \hat{y}_i)^2$
2-2.27 = -0.27	$(-0.27)^2 = 0.073$
4-3.43 = 0.57	$(0.57)^2 = 0.345$
0-(-0.05) = 0.05	$(0.05)^2 = 0.0025$
4-4.59 = -0.59	$(-0.59)^2 = 0.348$
6-5.75 = 0.25	$(0.25)^2 = 0.0625$
$\sum \hat{\varepsilon}_i = \sum (y_i - \hat{y}_i) = 0$	$\sum (y_i - \hat{y}_i)^2 = 0.811$

Example



$$\hat{y} = 1.11 + 1.16x$$

At $x = 3$:

$$y_i - \hat{y}_i = 4 - 4.59 = -.59$$

Logistic Regression

Logistic Regression

- Learn $P(Y|X)$ directly
- Consider learning $f: X \rightarrow Y$, where
 - X is a vector of real-valued features, $\langle X_1 \dots X_n \rangle$
 - Y is boolean
 - assume all X_i are conditionally independent given Y
 - model $P(X_i | Y = y_k)$ as Gaussian $N(\mu_{ik}, \sigma_i^2)$
 - model $P(Y)$ as Bernoulli (π)
 - Y is 1 with probability π

Derivation of $P(Y|X)$

$$P(Y=1|X) = \frac{P(Y=1)P(X|Y=1)}{P(Y=1)P(X|Y=1) + P(Y=0)P(X|Y=0)}$$

$$= \frac{1}{1 + \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)}}$$

$$= \frac{1}{1 + \exp(\ln \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)})}$$

$$= \frac{1}{1 + \exp((\ln \frac{1-\pi}{\pi}) + \sum_i \ln \frac{P(X_i|Y=0)}{P(X_i|Y=1)})}$$

$$P(x | y_k) = \frac{1}{\sigma_{ik}\sqrt{2\pi}} e^{\frac{-(x-\mu_{ik})^2}{2\sigma_{ik}^2}}$$

$$\sum_i \left(\frac{\mu_{i0} - \mu_{i1}}{\sigma_i^2} X_i + \frac{\mu_{i1}^2 - \mu_{i0}^2}{2\sigma_i^2} \right)$$

$$P(Y=1|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

Very Convenient

$$P(Y = 1|X = \langle X_1, \dots, X_n \rangle) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

implies

$$P(Y = 0|X = \langle X_1, \dots, X_n \rangle) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

implies

$$\frac{P(Y = 0|X)}{P(Y = 1|X)} = \exp(w_0 + \sum_i w_i X_i)$$

linear
classification
rule!

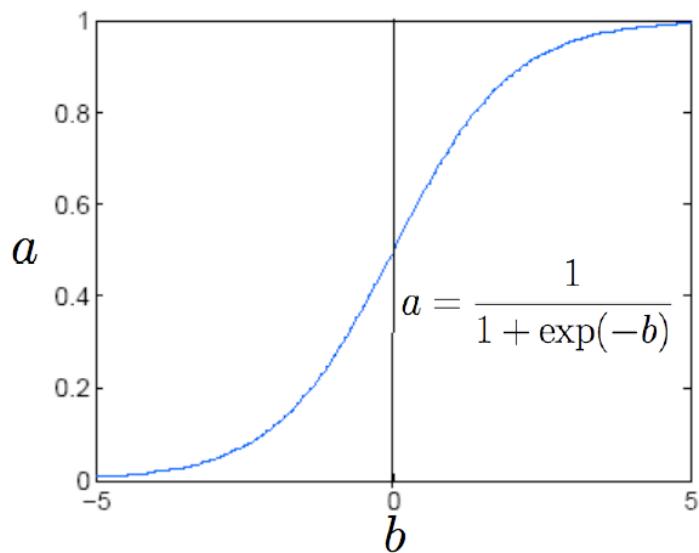
implies

$$\ln \frac{P(Y = 0|X)}{P(Y = 1|X)} = w_0 + \sum_i w_i X_i$$

Very Convenient

- Posteriors sum to 1 and remain in [0, 1]
- Logit: $l = \text{logit}(p) = \log\left(\frac{p}{1-p}\right) = \alpha + \beta x$
 - l is linear in x
- Probability: $p = \frac{e^l}{1+e^l}$

Logistic Function



- $p = 0, l = \log\left(\frac{p}{1-p}\right) = -\infty$
- $p = \frac{1}{2}, l = \log\left(\frac{p}{1-p}\right) = 0$
- $p = 1, l = \log\left(\frac{p}{1-p}\right) = +\infty$

$$P(Y = 1|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

Logistic Regression More Generally

- Logistic regression when Y is not boolean (but still discrete)

- $y \in \{y_1 \dots y_R\}$: learn $R-1$ sets of weights

- for $k < R$
$$P(Y = y_k | X) = \frac{\exp(w_{k0} + \sum_{i=1}^n w_{ki}X_i)}{1 + \sum_{j=1}^{R-1} \exp(w_{j0} + \sum_{i=1}^n w_{ji}X_i)}$$

- for $k=R$
$$P(Y = y_R | X) = \frac{1}{1 + \sum_{j=1}^{R-1} \exp(w_{j0} + \sum_{i=1}^n w_{ji}X_i)}$$

Training Logistic Regression: MCLE

- We have L training examples: $\{\langle X^1, Y^1 \rangle, \dots \langle X^L, Y^L \rangle\}$
- Maximum likelihood estimate for parameters W

$$\begin{aligned} W_{MLE} &= \arg \max_W P(\langle X^1, Y^1 \rangle \dots \langle X^L, Y^L \rangle | W) \\ &= \arg \max_W \prod_l P(\langle X^l, Y^l \rangle | W) \end{aligned}$$

- Maximum Conditional Likelihood Estimate (MCLE)

Training Logistic Regression: MCLE

- Choose parameters $\langle w_0 \dots w_n \rangle$ to maximize conditional likelihood of training data

$$P(Y = 0|X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$P(Y = 1|X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

- Training data $D = \{\langle X^1, Y^1 \rangle, \dots \langle X^L, Y^L \rangle\}$
- Data likelihood = $\prod_l P(X^l, Y^l | W)$
- Data conditional likelihood = $\prod_l P(Y^l | X^l, W)$

$$W_{MCLE} = \arg \max_W \prod_l P(Y^l | W, X^l)$$

Conditional Log Likelihood

$$l(W) \equiv \ln \prod_l P(Y^l | X^l, W) = \sum_l \ln P(Y^l | X^l, W)$$

$$P(Y = 0 | X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$P(Y = 1 | X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$\begin{aligned} l(W) &= \sum_l Y^l \ln P(Y^l = 1 | X^l, W) + (1 - Y^l) \ln P(Y^l = 0 | X^l, W) \\ &= \sum_l Y^l \ln \frac{P(Y^l = 1 | X^l, W)}{P(Y^l = 0 | X^l, W)} + \ln P(Y^l = 0 | X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i w_i X_i^l)) \end{aligned}$$

Maximizing Conditional Log Likelihood

$$P(Y = 0|X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

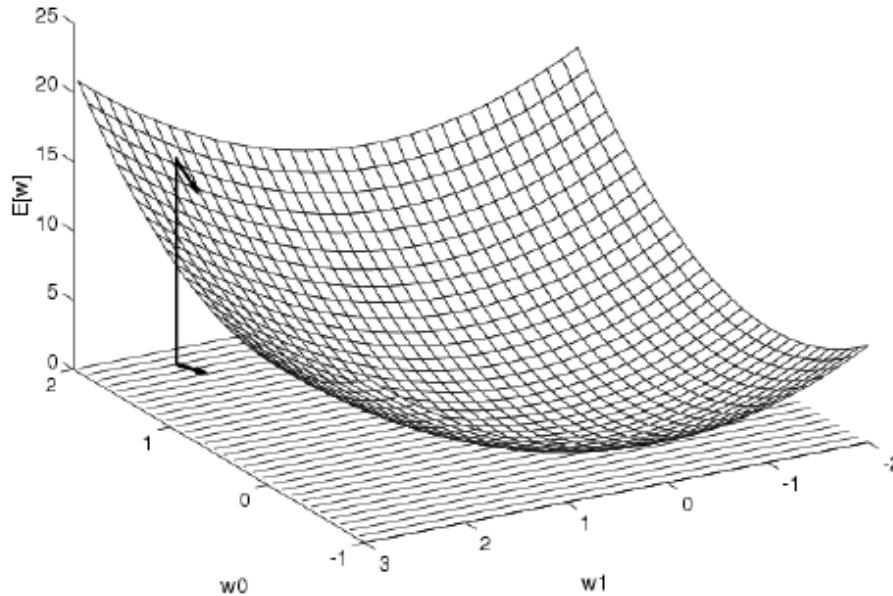
$$P(Y = 1|X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$\begin{aligned} l(W) &\equiv \ln \prod_l P(Y^l | X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i^n w_i X_i^l)) \end{aligned}$$

Good news: $l(W)$ is concave function of W

Bad news: no closed-form solution to maximize $l(W)$

Gradient Descent



Gradient

$$\nabla E[\vec{w}] \equiv \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$

Training rule:

$$\Delta \vec{w} = -\eta \nabla E[\vec{w}]$$

i.e.,

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

Maximize Conditional Log Likelihood: Gradient Ascent

$$\begin{aligned} l(W) &\equiv \ln \prod_l P(Y^l | X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i w_i X_i^l)) \end{aligned}$$

$$\frac{\partial l(W)}{\partial w_i} = \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

Maximize Conditional Log Likelihood: Gradient Ascent

$$\begin{aligned} l(W) &\equiv \ln \prod_l P(Y^l | X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i^n w_i X_i^l)) \end{aligned}$$

$$\frac{\partial l(W)}{\partial w_i} = \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

Gradient ascent algorithm: iterate until change $< \varepsilon$

For all i , repeat

$$w_i \leftarrow w_i + \eta \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

Logistic Regression: Summary

- Consider learning $f: X \rightarrow Y$, where
 - X is a vector of real-valued features, $\langle X_1 \dots X_n \rangle$
 - Y is boolean
 - assume all X_i are conditionally independent given Y
 - model $P(X_i | Y = y_k)$ as Gaussian $N(\mu_{ik}, \sigma_i^2)$
 - model $P(Y)$ as Bernoulli (π)
- Then $P(Y|X)$ is of this form and we can directly estimate W

$$P(Y = 1 | X = \langle X_1, \dots, X_n \rangle) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

Linear Discriminant Functions

Augmented Feature Vector

- Linear discriminant function: $g(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$
- Can rewrite it:

$$g(\mathbf{x}) = \underbrace{[\mathbf{w}_0 \quad \mathbf{w}^t]}_{\text{new weight vector } \mathbf{a}} \underbrace{\begin{bmatrix} 1 \\ \mathbf{x} \end{bmatrix}}_{\text{new feature vector } \mathbf{y}} = \mathbf{a}^t \mathbf{y} = g(\mathbf{y})$$

- \mathbf{y} is called the **augmented feature vector**
- Added a dummy dimension to get a completely equivalent new **homogeneous problem**

old problem

$$g(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

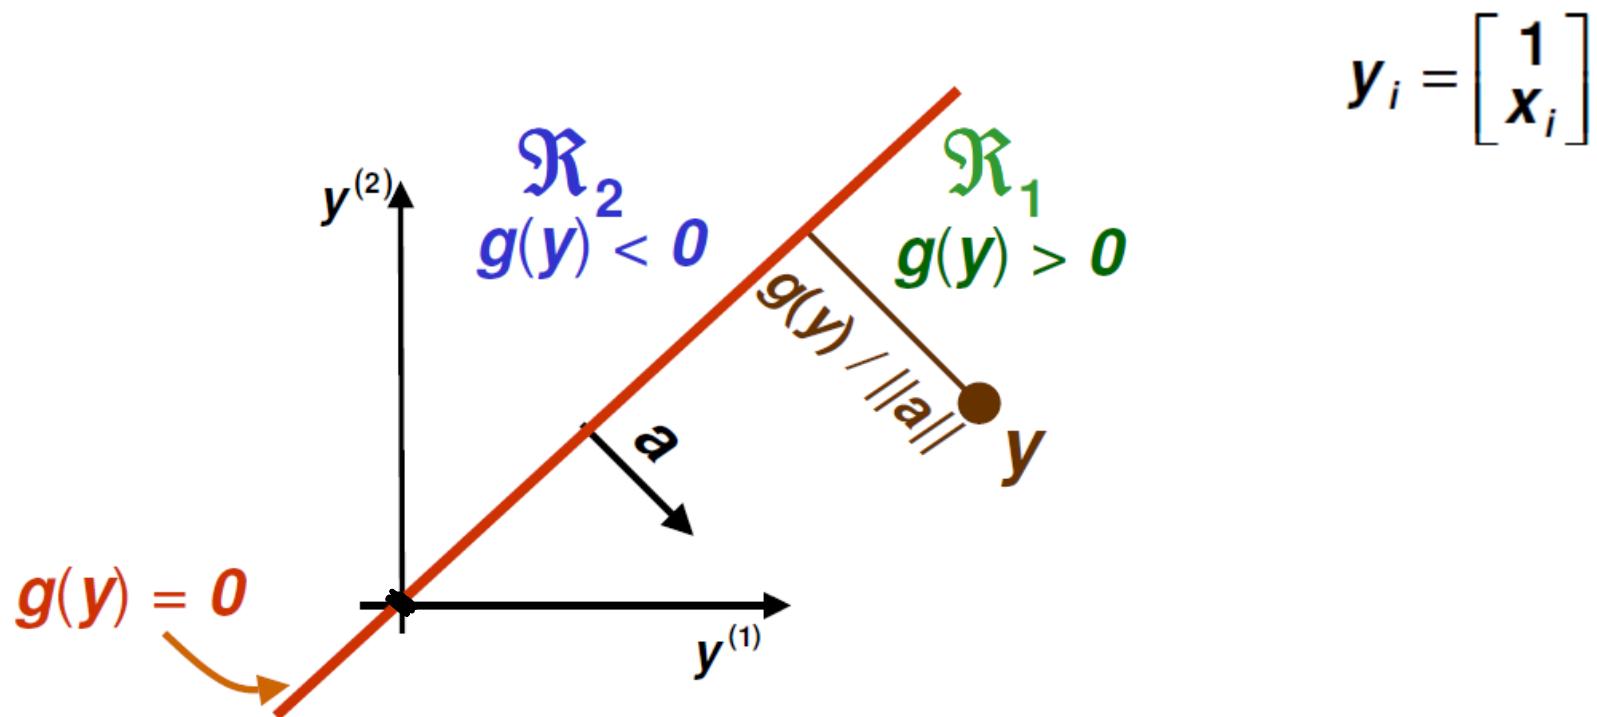
$$\begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_d \end{bmatrix}$$

new problem

$$g(\mathbf{y}) = \mathbf{a}^t \mathbf{y}$$

$$\begin{bmatrix} 1 \\ \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_d \end{bmatrix}$$

- Feature augmentation is done for simpler notation
- From now on, always assume that we have augmented feature vectors
 - Given samples $\mathbf{x}_1, \dots, \mathbf{x}_n$ convert them to augmented samples $\mathbf{y}_1, \dots, \mathbf{y}_n$ by adding a new dimension of value 1



Training Error

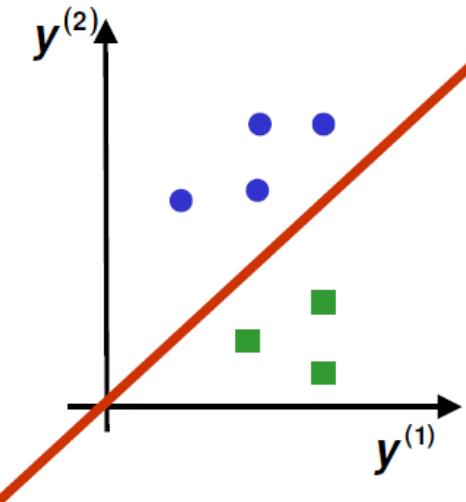
- For the rest of this part, assume we have 2 classes
 - Samples: $\mathbf{y}_1, \dots, \mathbf{y}_n$, some in class 1, some in class 2
- Use samples to determine weights \mathbf{a} in the discriminant function $\mathbf{g}(\mathbf{y}) = \mathbf{a}^t \mathbf{y}$
- What should the criterion for determining \mathbf{a} be?
- For now, suppose we want to minimize the training error (the number of misclassified samples $\mathbf{y}_1, \dots, \mathbf{y}_n$)
- Recall that:
 - $\mathbf{g}(\mathbf{y}_i) > 0 \Rightarrow \mathbf{y}_i$ classified as c_1
 - $\mathbf{g}(\mathbf{y}_i) < 0 \Rightarrow \mathbf{y}_i$ classified as c_2
- Thus training error is **0** if
$$\begin{cases} \mathbf{g}(\mathbf{y}_i) > 0 & \forall \mathbf{y}_i \in c_1 \\ \mathbf{g}(\mathbf{y}_i) < 0 & \forall \mathbf{y}_i \in c_2 \end{cases}$$

“Normalization”

- Thus training error is 0 if:
$$\begin{cases} \mathbf{a}^t \mathbf{y}_i > 0 & \forall \mathbf{y}_i \in \mathcal{C}_1 \\ \mathbf{a}^t \mathbf{y}_i < 0 & \forall \mathbf{y}_i \in \mathcal{C}_2 \end{cases}$$
- Equivalently, training error is 0 if:
$$\begin{cases} \mathbf{a}^t \mathbf{y}_i > 0 & \forall \mathbf{y}_i \in \mathcal{C}_1 \\ \mathbf{a}^t (-\mathbf{y}_i) > 0 & \forall \mathbf{y}_i \in \mathcal{C}_2 \end{cases}$$
- This suggests “normalization” (a.k.a. reflection):
 1. Replace all examples from class 2 by:
$$\mathbf{y}_i \rightarrow -\mathbf{y}_i \quad \forall \mathbf{y}_i \in \mathcal{C}_2$$
 2. Seek weight vector \mathbf{a} such that
$$\mathbf{a}^t \mathbf{y}_i > 0 \quad \forall \mathbf{y}_i$$
 - If such \mathbf{a} exists, it is called a separating or solution vector
 - Original samples $\mathbf{x}_1, \dots, \mathbf{x}_n$ can indeed be separated by a line

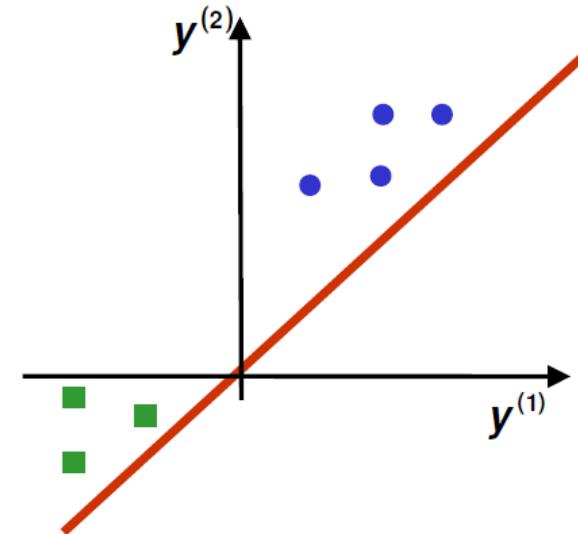
Normalization

before normalization



- Seek a hyperplane that separates patterns from different categories

after “normalization”

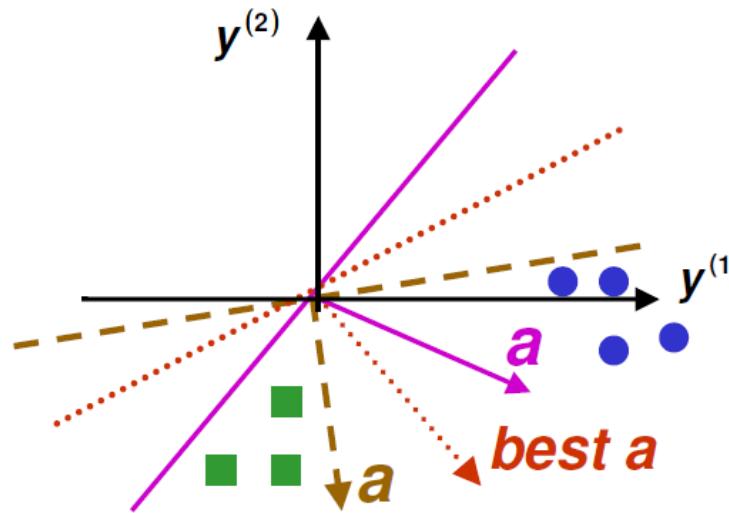


- Seek hyperplane that puts *normalized* patterns on the same (positive) side

Solution Region

- Find weight vector \mathbf{a} such that for all samples:

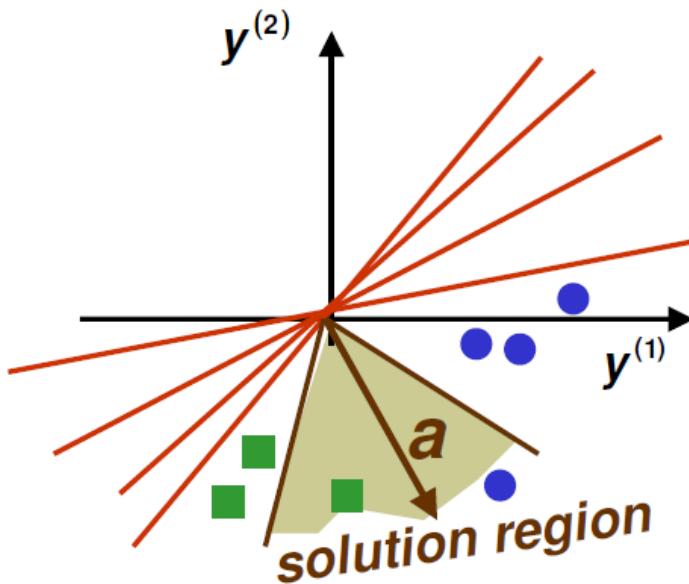
$$\mathbf{a}^t \mathbf{y}_i = \sum_{k=0}^d \mathbf{a}_k y_i^{(k)} > 0$$



- In general, there can be many solutions

Solution Region

- **Solution region** for \mathbf{a} : set of all possible solutions defined in terms of normal \mathbf{a} to the separating hyperplane



Optimization

- Need to minimize a function of many variables

$$J(\mathbf{x}) = J(x_1, \dots, x_d)$$

- We know how to minimize $J(\mathbf{x})$
 - Take partial derivatives and set them to zero

$$\begin{bmatrix} \frac{\partial}{\partial x_1} J(\mathbf{x}) \\ \vdots \\ \frac{\partial}{\partial x_d} J(\mathbf{x}) \end{bmatrix} = \nabla J(\mathbf{x}) = \mathbf{0}$$

gradient 

Optimization

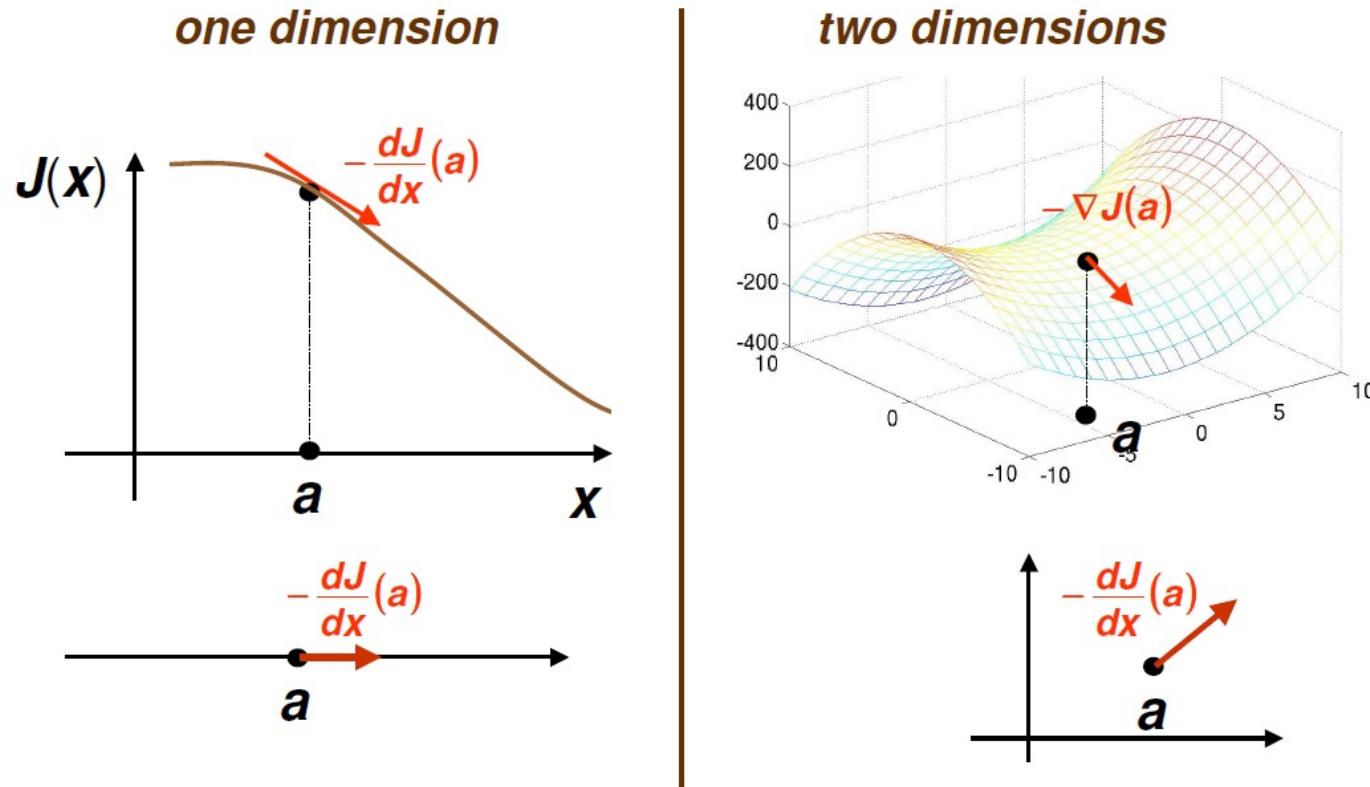
- However solving analytically is not always easy
 - For example:

$$\begin{cases} \sin(x_1^2 + x_2^3) + e^{x_4} = 0 \\ \cos(x_1^2 + x_2^3) + \log(x_5^{x_4}) = 0 \end{cases}$$

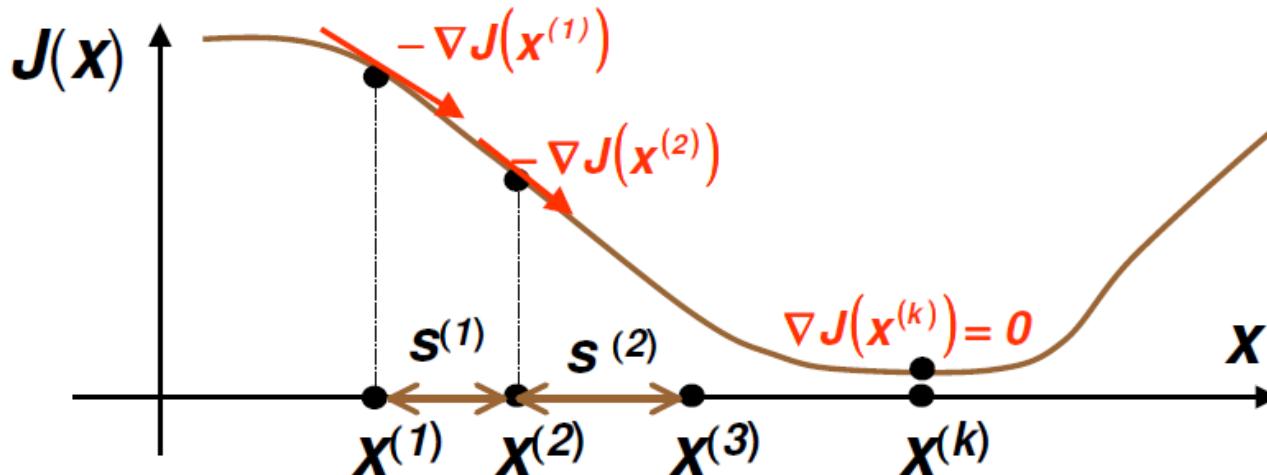
- Sometimes it is not even possible to write down an analytical expression for the derivative (example later today)

Gradient Descent

- Gradient $\nabla J(x)$ points in direction of steepest increase of $J(x)$, and $-\nabla J(x)$ in direction of steepest decrease



Gradient Descent



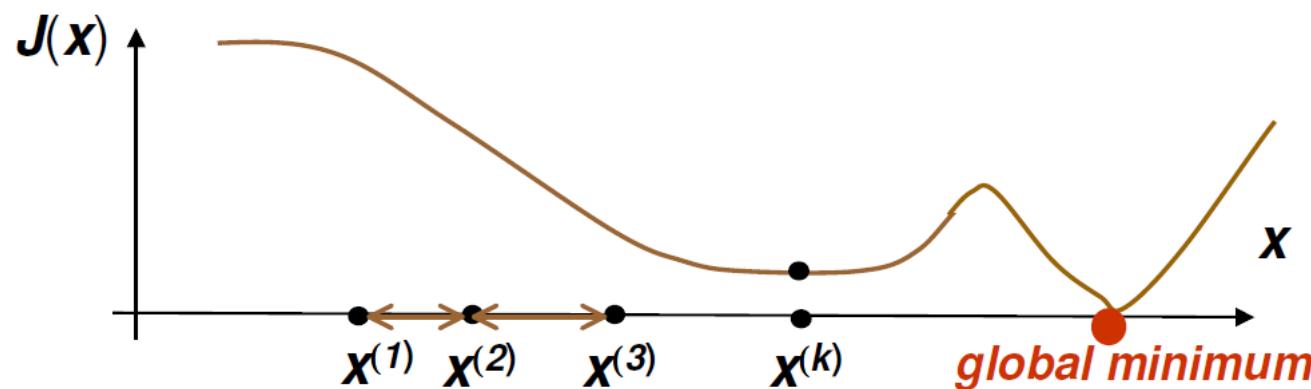
Gradient Descent for minimizing any function $J(x)$

- Set $k = 1$ and $x^{(1)}$ to some initial guess for the weight vector
- While $\eta^{(k)} |\nabla J(x^{(k)})| > \epsilon$
 - Choose learning rate $\eta^{(k)}$

$$\begin{aligned} \mathbf{x}^{(k+1)} &= \mathbf{x}^{(k)} - \eta^{(k)} \nabla J(\mathbf{x}) \\ k &= k + 1 \end{aligned} \quad (\text{update rule})$$

Gradient Descent

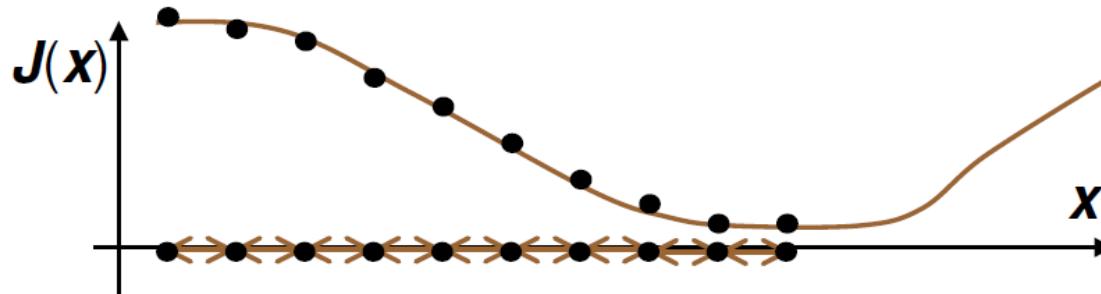
- Gradient decent is guaranteed to only find **local minima**



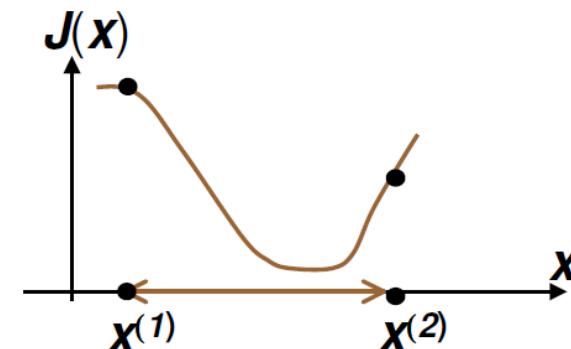
- Nevertheless gradient descent is very popular because it is simple and applicable to any function

Gradient Descent

- Main issue: how to set parameter η (learning rate)
 - If η is too small, too many iterations



- If η is too large may overshoot the minimum and possibly never find it



LDF Criterion Function

- Find weight vector \mathbf{a} such that for all samples $\mathbf{y}_1, \dots, \mathbf{y}_n$

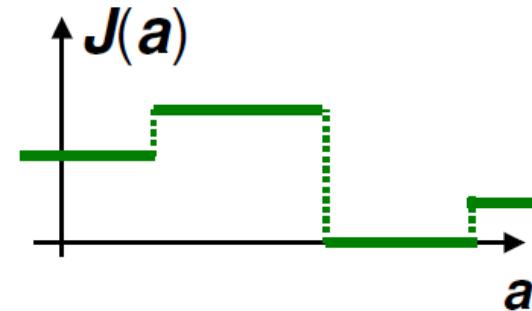
$$\mathbf{a}^t \mathbf{y}_i = \sum_{k=0}^d a_k y_i^{(k)} > 0$$

- Need criterion function $J(\mathbf{a})$ which is minimized when \mathbf{a} is a solution vector
- Let \mathcal{Y}_M be the set of examples misclassified by \mathbf{a}

$$\mathcal{Y}_M(\mathbf{a}) = \{ \mathbf{y}_i \text{ s.t. } \mathbf{a}^t \mathbf{y}_i < 0 \}$$

- First natural choice: number of misclassified examples

$$J(\mathbf{a}) = |\mathcal{Y}_M(\mathbf{a})|$$



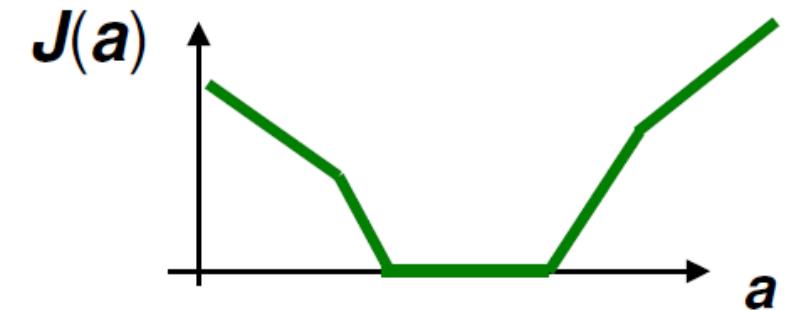
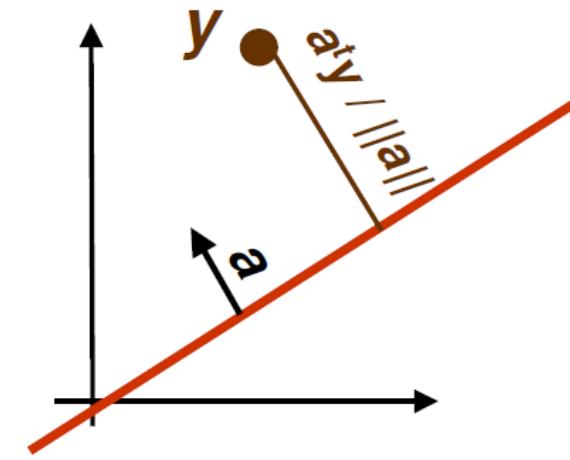
- Piecewise constant, gradient descent is useless

Perceptron

Perceptron Criterion Function

$$J_p(a) = \sum_{y \in Y_M} (-a^t y)$$

- If y is misclassified, $a^t y < 0$
- Thus $J_p(a) > 0$
- $J_p(a)$ is $\|a\|$ times the sum of distances of misclassified examples to decision boundary
- $J_p(a)$ is piecewise linear and thus suitable for gradient descent



Perceptron Batch Rule

$$J_p(\mathbf{a}) = \sum_{y \in Y_M} (-\mathbf{a}^t \mathbf{y})$$

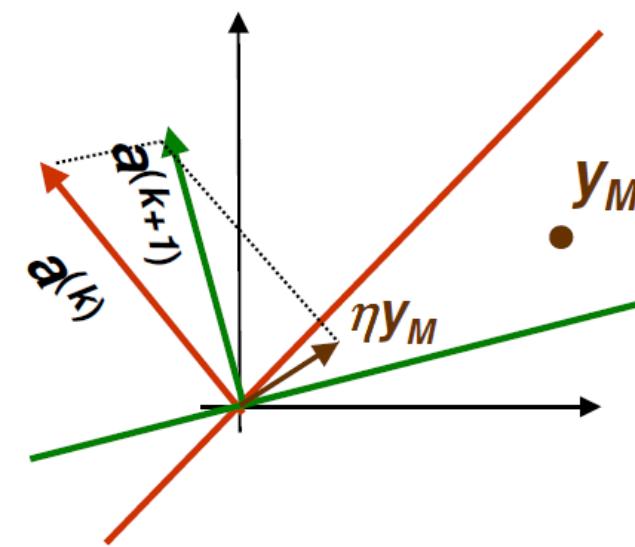
- Gradient of $J_p(\mathbf{a})$ is: $\nabla J_p(\mathbf{a}) = \sum_{y \in Y_M} (-\mathbf{y})$
 - Y_M are samples misclassified by $\mathbf{a}^{(k)}$
 - It is not possible to solve $\nabla J_p(\mathbf{a}) = \mathbf{0}$ analytically because of Y_M
- Update rule for gradient descent: $\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} - \eta^{(k)} \nabla J(x)$
- Thus the *gradient decent batch update rule for $J_p(\mathbf{a})$* is:
$$\mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \eta^{(k)} \sum_{y \in Y_M} \mathbf{y}$$
- It is called batch rule because it is based on all misclassified examples

Perceptron Single Sample Rule

- The *gradient decent single sample rule* for $J_p(a)$ is:

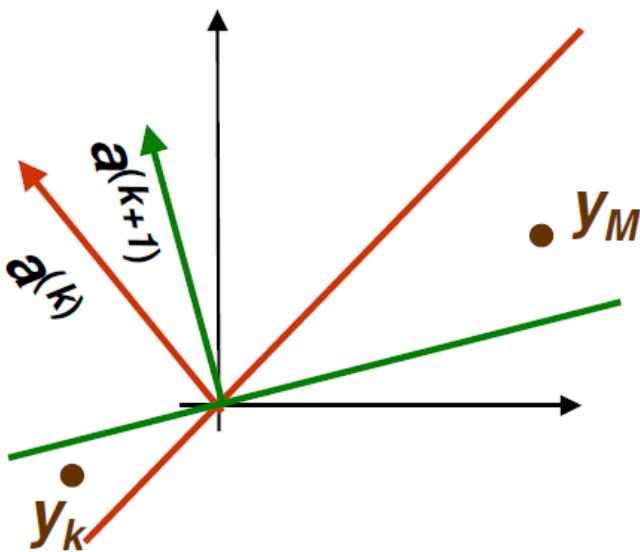
$$a^{(k+1)} = a^{(k)} + \eta^{(k)} y_M$$

- Note that y_M is one sample misclassified by $a(k)$
- Must have a consistent way of visiting samples
- Geometric Interpretation:
 - y_M misclassified by $a^{(k)}$ $(a^{(k)})^t y_M \leq 0$
 - y_M is on the wrong side of decision hyperplane
 - Adding ηy_M to a moves the new decision hyperplane in the right direction with respect to y_M

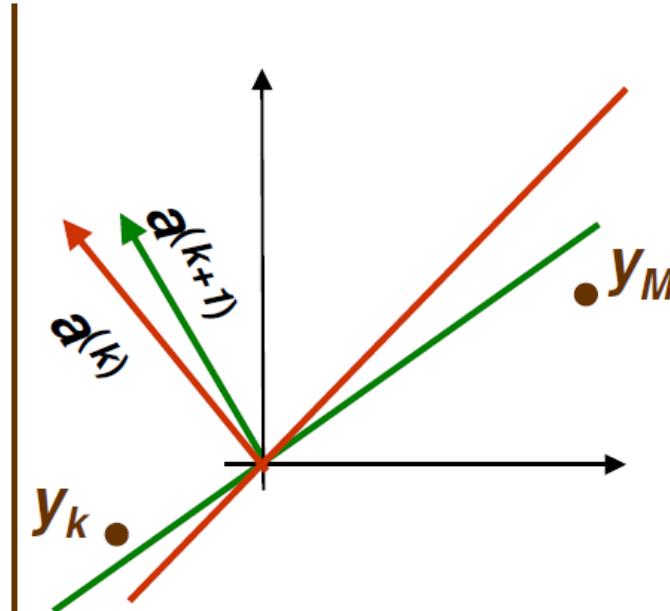


Perceptron Single Sample Rule

$$\mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \eta^{(k)} y_M$$



η is too large, previously
correctly classified sample
 y_k is now misclassified



η is too small, y_M is still
misclassified

Perceptron Example

	features					grade
<i>name</i>	<i>good attendance?</i>	<i>tall?</i>	<i>sleeps in class?</i>	<i>chews gum?</i>		
Jane	yes (1)	yes (1)	no (-1)	no (-1)	A	
Steve	yes (1)	yes (1)	yes (1)	yes (1)	F	
Mary	no (-1)	no (-1)	no (-1)	yes (1)	F	
Peter	yes (1)	no (-1)	no (-1)	yes (1)	A	

- Class 1: students who get A
- Class 2: students who get F

	features					grade
<i>name</i>	<i>extra</i>	<i>good attendance?</i>	<i>tall?</i>	<i>sleeps in class?</i>	<i>chews gum?</i>	
Jane	1	yes (1)	yes (1)	no (-1)	no (-1)	A
Steve	1	yes (1)	yes (1)	yes (1)	yes (1)	F
Mary	1	no (-1)	no (-1)	no (-1)	yes (1)	F
Peter	1	yes (1)	no (-1)	no (-1)	yes (1)	A

- Augment samples by adding an extra feature (dimension) equal to 1

	features					grade
<i>name</i>	<i>extra</i>	<i>good attendance?</i>	<i>tall?</i>	<i>sleeps in class?</i>	<i>chews gum?</i>	
Jane	1	yes (1)	yes (1)	no (-1)	no (-1)	A
Steve	-1	yes (-1)	yes (-1)	yes (-1)	yes (-1)	F
Mary	-1	no (1)	no (1)	no (1)	yes (-1)	F
Peter	1	yes (1)	no (-1)	no (-1)	yes (1)	A

- Normalize:
 - Replace all examples from class 2 by their negative values
 - Seek \mathbf{a} such that:

$$\mathbf{y}_i \rightarrow -\mathbf{y}_i \quad \forall \mathbf{y}_i \in \mathcal{C}_2$$

$$\mathbf{a}^t \mathbf{y}_i > 0 \quad \forall \mathbf{y}_i$$

	features					grade
<i>name</i>	<i>extra</i>	<i>good attendance?</i>	<i>tall?</i>	<i>sleeps in class?</i>	<i>chews gum?</i>	
Jane	1	yes (1)	yes (1)	no (-1)	no (-1)	A
Steve	-1	yes (-1)	yes (-1)	yes (-1)	yes (-1)	F
Mary	-1	no (1)	no (1)	no (1)	yes (-1)	F
Peter	1	yes (1)	no (-1)	no (-1)	yes (1)	A

- Single Sample Rule

- Sample is misclassified if $\mathbf{a}^t \mathbf{y}_i = \sum_{k=0}^4 a_k y_i^{(k)} < 0$
- Gradient descent single sample rule: $\mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \eta^{(k)} \sum_{y \in Y_M} y$
- Set η fixed learning rate to $\eta^{(k)} = 1$: $\boxed{\mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \mathbf{y}_M}$

- Set equal initial weights
 $\mathbf{a}^{(1)} = [0.25, 0.25, 0.25, 0.25, 0.25]$
- Visit all samples sequentially, modifying the weights after each misclassified example

<i>name</i>	$\mathbf{a}^t \mathbf{y}$	<i>misclassified?</i>
Jane	$0.25*1+0.25*1+0.25*1+0.25*(-1)+0.25*(-1) > 0$	<i>no</i>
Steve	$0.25*(-1)+0.25*(-1)+0.25*(-1)+0.25*(-1)+0.25*(-1) < 0$	<i>yes</i>

- New weights

$$\begin{aligned}
 \mathbf{a}^{(2)} &= \mathbf{a}^{(1)} + \mathbf{y}_M = [0.25 \ 0.25 \ 0.25 \ 0.25 \ 0.25] + \\
 &\quad + [-1 \ -1 \ -1 \ -1 \ -1] = \\
 &= [-0.75 \ -0.75 \ -0.75 \ -0.75 \ -0.75]
 \end{aligned}$$

$$\mathbf{a}^{(2)} = [-0.75 \ -0.75 \ -0.75 \ -0.75 \ -0.75]$$

<i>name</i>	$\mathbf{a}^t \mathbf{y}$	<i>misclassified?</i>
Mary	$-0.75*(-1) - 0.75*1 - 0.75 * 1 - 0.75 * 1 - 0.75*(-1) < 0$	yes

- New weights

$$\begin{aligned}
 \mathbf{a}^{(3)} &= \mathbf{a}^{(2)} + \mathbf{y}_M = [-0.75 \ -0.75 \ -0.75 \ -0.75 \ -0.75] + \\
 &\quad + [-1 \ 1 \ 1 \ 1 \ -1] = \\
 &= [-1.75 \ 0.25 \ 0.25 \ 0.25 \ -1.75]
 \end{aligned}$$

$$\mathbf{a}^{(3)} = [-1.75 \ 0.25 \ 0.25 \ 0.25 \ -1.75]$$

<i>name</i>	$\mathbf{a}^t \mathbf{y}$	<i>misclassified?</i>
Peter	$-1.75 * 1 + 0.25 * 1 + 0.25 * (-1) + 0.25 * (-1) - 1.75 * 1 < 0$	yes

- New weights

$$\begin{aligned}
 \mathbf{a}^{(4)} &= \mathbf{a}^{(3)} + \mathbf{y}_M = [-1.75 \ 0.25 \ 0.25 \ 0.25 \ -1.75] + \\
 &\quad + [1 \ 1 \ -1 \ -1 \ 1] = \\
 &= [-0.75 \ 1.25 \ -0.75 \ -0.75 \ -0.75]
 \end{aligned}$$

$$\mathbf{a}^{(4)} = [-0.75 \ 1.25 \ -0.75 \ -0.75 \ -0.75]$$

<i>name</i>	$\mathbf{a}^t \mathbf{y}$	<i>misclassified?</i>
Jane	$-0.75 * 1 + 1.25 * 1 - 0.75 * 1 - 0.75 * (-1) - 0.75 * (-1) + 0$	<i>no</i>
Steve	$-0.75 * (-1) + 1.25 * (-1) - 0.75 * (-1) - 0.75 * (-1) - 0.75 * (-1) > 0$	<i>no</i>
Mary	$-0.75 * (-1) + 1.25 * 1 - 0.75 * 1 - 0.75 * 1 - 0.75 * (-1) > 0$	<i>no</i>
Peter	$-0.75 * 1 + 1.25 * 1 - 0.75 * (-1) - 0.75 * (-1) - 0.75 * 1 > 0$	<i>no</i>

- Thus the discriminant function is:

$$g(\mathbf{y}) = -0.75 * y^{(0)} + 1.25 * y^{(1)} - 0.75 * y^{(2)} - 0.75 * y^{(3)} - 0.75 * y^{(4)}$$

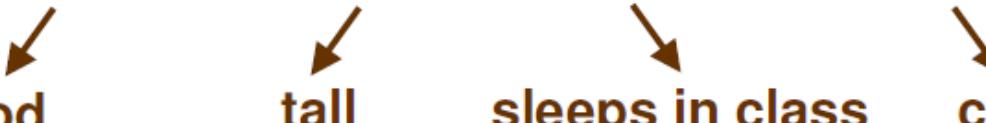
- Converting back to the original features \mathbf{x} :

$$g(\mathbf{x}) = 1.25 * x^{(1)} - 0.75 * x^{(2)} - 0.75 * x^{(3)} - 0.75 * x^{(4)} - 0.75$$

- Converting back to the original features x :

$$1.25 * x^{(1)} - 0.75 * x^{(2)} - 0.75 * x^{(3)} - 0.75 * x^{(4)} > 0.75 \Rightarrow grade\ A$$

$$1.25 * x^{(1)} - 0.75 * x^{(2)} - 0.75 * x^{(3)} - 0.75 * x^{(4)} < 0.75 \Rightarrow grade\ F$$



 good attendance tall sleeps in class chews gum

- This is just one possible solution vector
- If we started with weights

$$a^{(1)} = [0, 0.5, 0.5, 0, 0],$$

- The solution would be $[-1, 1.5, -0.5, -1, -1]$

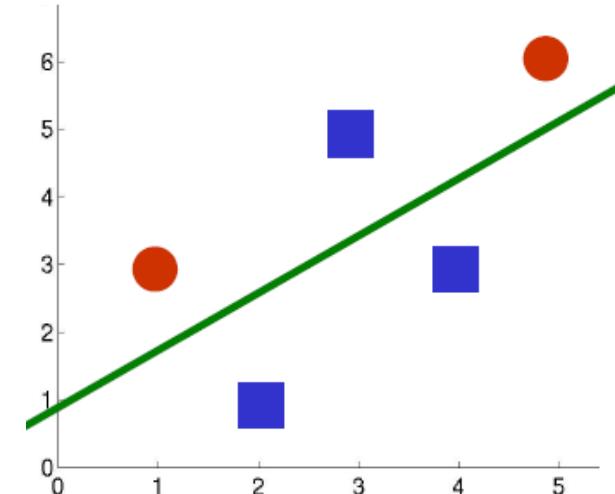
$$1.5 * x^{(1)} - 0.5 * x^{(2)} - x^{(3)} - x^{(4)} > 1 \Rightarrow grade\ A$$

$$1.5 * x^{(1)} - 0.5 * x^{(2)} - x^{(3)} - x^{(4)} < 1 \Rightarrow grade\ F$$

- In this solution, being tall is the least important feature

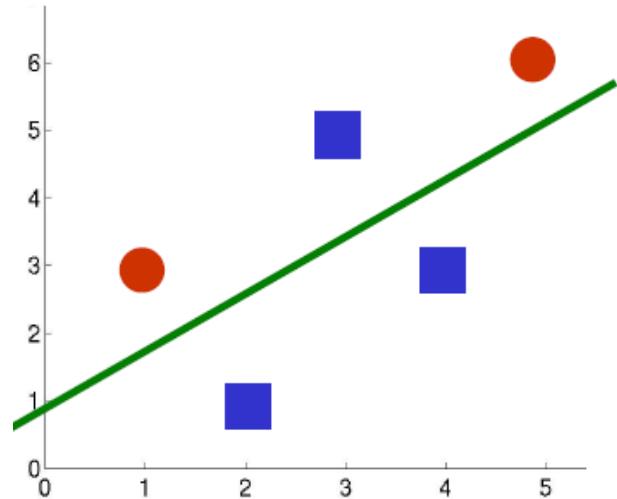
LDF: Non-separable Example

- Suppose we have 2 features and the samples are:
 - Class 1: [2,1], [4,3], [3,5]
 - Class 2: [1,3] and [5,6]
- These samples are not separable by a line
- Still would like to get approximate separation by a line
 - A good choice is shown in green
 - Some samples may be “noisy”, and we could accept them being misclassified



LDF: Non-separable Example

- Obtain y_1, y_2, y_3, y_4 by adding extra feature and “normalizing”



$$y_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad y_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad y_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad y_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad y_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

LDF: Non-separable Example

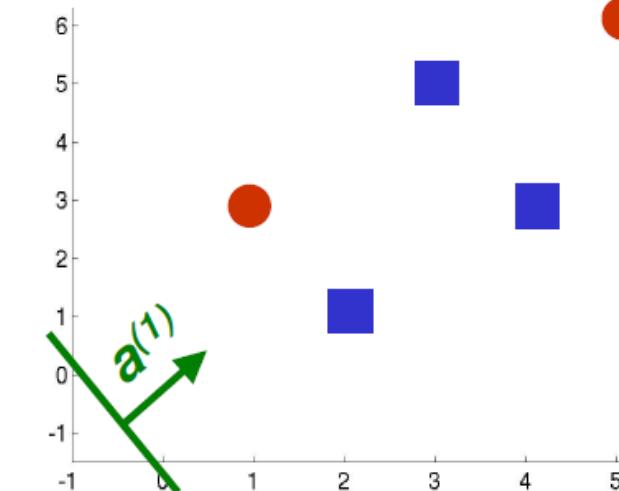
- Apply Perceptron single sample algorithm

- Initial equal weights

$$\mathbf{a}^{(1)} = [1 \ 1 \ 1]$$

- Line equation $x^{(1)}+x^{(2)}+1=0$

- Fixed learning rate $\eta = 1$



$$\mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \mathbf{y}_M$$

$$\mathbf{y}_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad \mathbf{y}_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad \mathbf{y}_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad \mathbf{y}_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad \mathbf{y}_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

- $\mathbf{y}_1^t \mathbf{a}^{(1)} = [1 \ 1 \ 1]^* [1 \ 2 \ 1]^t > 0 \quad \checkmark$
- $\mathbf{y}_2^t \mathbf{a}^{(1)} = [1 \ 1 \ 1]^* [1 \ 4 \ 3]^t > 0 \quad \checkmark$
- $\mathbf{y}_3^t \mathbf{a}^{(1)} = [1 \ 1 \ 1]^* [1 \ 3 \ 5]^t > 0 \quad \checkmark$

LDF: Non-separable Example

$$\mathbf{a}^{(1)} = [1 \ 1 \ 1] \quad \mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \mathbf{y}_M$$

$$\mathbf{y}_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad \mathbf{y}_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad \mathbf{y}_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad \mathbf{y}_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad \mathbf{y}_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

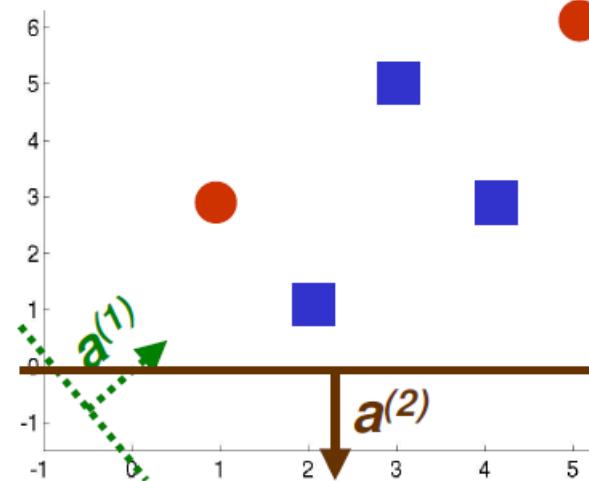
- $\mathbf{y}_4^T \mathbf{a}^{(1)} = [1 \ 1 \ 1]^* [-1 \ -1 \ -3]^t = -5 < 0$

$$\mathbf{a}^{(2)} = \mathbf{a}^{(1)} + \mathbf{y}_M = [1 \ 1 \ 1] + [-1 \ -1 \ -3] = [0 \ 0 \ -2]$$

- $\mathbf{y}_5^T \mathbf{a}^{(2)} = [0 \ 0 \ -2]^* [-1 \ -5 \ -6]^t = 12 > 0 \quad \checkmark$

- $\mathbf{y}_1^T \mathbf{a}^{(2)} = [0 \ 0 \ -2]^* [1 \ 2 \ 1]^t < 0$

$$\mathbf{a}^{(3)} = \mathbf{a}^{(2)} + \mathbf{y}_M = [0 \ 0 \ -2] + [1 \ 2 \ 1] = [1 \ 2 \ -1]$$



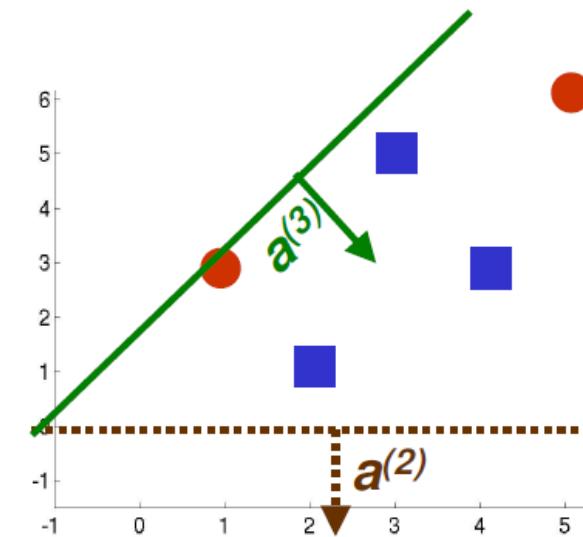
LDF: Non-separable Example

$$\mathbf{a}^{(3)} = [1 \ 2 \ -1] \quad \mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \mathbf{y}_M$$

$$\mathbf{y}_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad \mathbf{y}_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad \mathbf{y}_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad \mathbf{y}_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad \mathbf{y}_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

- $\mathbf{y}_2^T \mathbf{a}^{(3)} = [1 \ 4 \ 3]^* [1 \ 2 \ -1]^t = 6 > 0$ ✓
- $\mathbf{y}_3^T \mathbf{a}^{(3)} = [1 \ 3 \ 5]^* [1 \ 2 \ -1]^t > 0$ ✓
- $\mathbf{y}_4^T \mathbf{a}^{(3)} = [-1 \ -1 \ -3]^* [1 \ 2 \ -1]^t = 0$

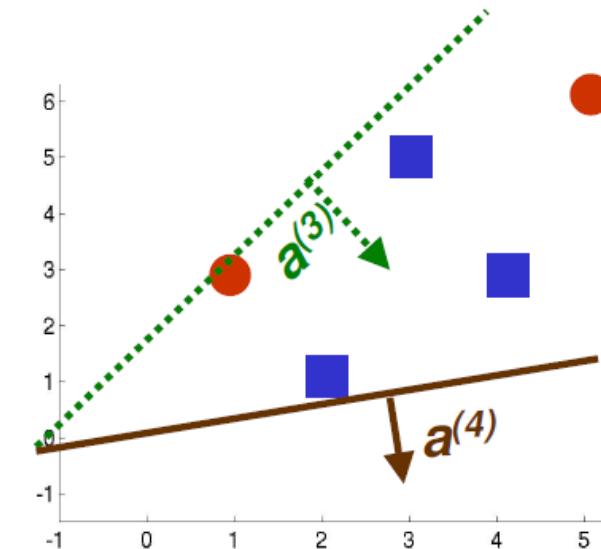
$$\mathbf{a}^{(4)} = \mathbf{a}^{(3)} + \mathbf{y}_M = [1 \ 2 \ -1] + [-1 \ -1 \ -3] = [0 \ 1 \ -4]$$



LDF: Non-separable Example

$$\mathbf{a}^{(4)} = [0 \ 1 \ -4] \quad \mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \mathbf{y}_M$$

$$\mathbf{y}_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad \mathbf{y}_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad \mathbf{y}_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad \mathbf{y}_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad \mathbf{y}_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$



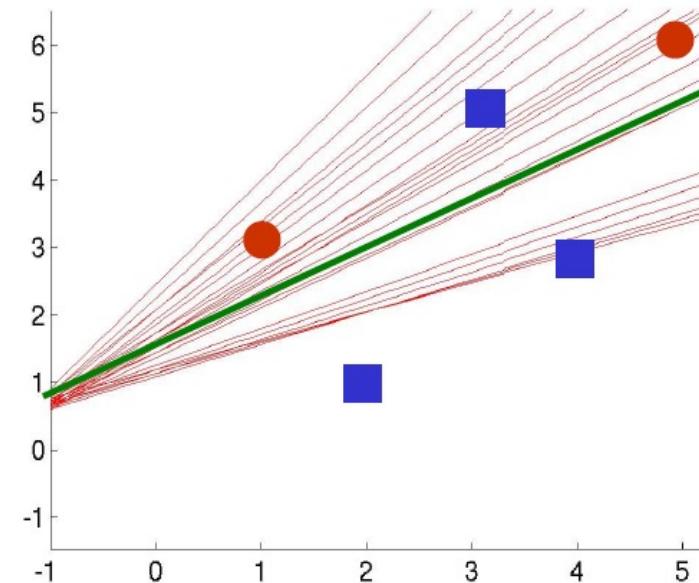
- $\mathbf{y}_5^T \mathbf{a}^{(4)} = [-1 \ -5 \ -6] * [0 \ 1 \ -4] = 19 > 0$
- $\mathbf{y}_1^T \mathbf{a}^{(4)} = [1 \ 2 \ 1] * [0 \ 1 \ -4] = -2 < 0$
-

LDF: Non-separable Example

- We can continue this forever
- There is no solution vector \mathbf{a} satisfying for all i

$$\mathbf{a}^t \mathbf{y}_i = \sum_{k=0}^5 a_k y_i^{(k)} > 0$$

- Need to stop but at a good point
- Solutions at iterations 900 through 915
 - Some are good and some are not
- How do we stop at a good solution?



Convergence of Perceptron Rules

- If classes are linearly separable and we use fixed learning rate, that is for $\eta^{(k)} = \text{const}$
- Then, *both the single sample and batch perceptron rules converge to a correct solution (could be any a in the solution space)*
- If classes are not linearly separable:
 - The algorithm does not stop, it keeps looking for a solution which does not exist

Convergence of Perceptron Rules

- If classes are not linearly separable:
 - By choosing appropriate learning rate, we can always ensure convergence: $\eta^{(k)} \rightarrow 0$ as $k \rightarrow \infty$
 - For example inverse linear learning rate: $\eta^{(k)} = \frac{\eta^{(1)}}{k}$
 - For inverse linear learning rate, convergence in the linearly separable case can also be proven
 - **No guarantee that we stopped at a good point**, but there are good reasons to choose inverse linear learning rate