

Machine Learning: Models and Applications

Lecture 8

Lecturer: Xinchao Wang

xinchao@nus.edu.sg

National University of Singapore

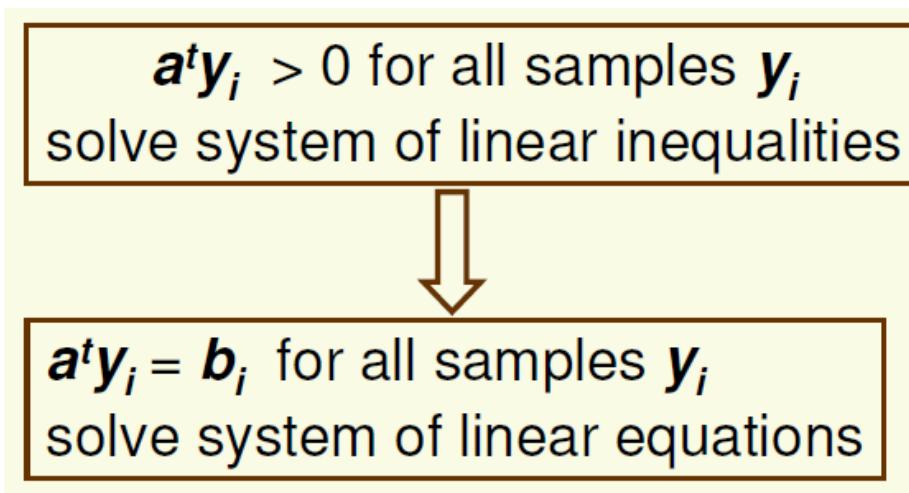
Overview

- Minimum Squared Error (MSE)
- Support Vector Machines (SVM)
 - Introduction
 - Linear Discriminant
 - Linearly Separable Case
 - Linearly Non Separable Case
 - Kernel Trick
 - Non Linear Discriminant
 - Multi-class SVMs

Minimum Squared-Error Procedures

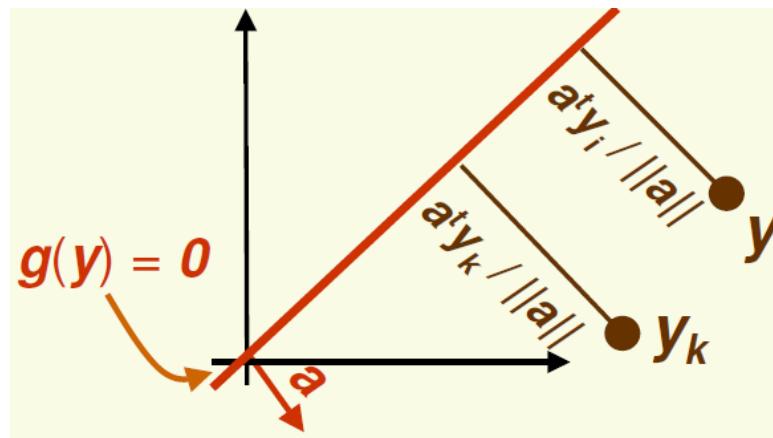
Minimum Squared-Error Procedures

- Idea: convert to easier and better understood problem



- MSE procedure
 - Choose positive constants $\mathbf{b}_1, \mathbf{b}_2, \dots, \mathbf{b}_n$
 - Try to find weight vector \mathbf{a} such that $\mathbf{a}^t \mathbf{y}_i = \mathbf{b}_i$ for all samples \mathbf{y}_i
 - If we can find such a vector, then \mathbf{a} is a solution because the \mathbf{b}_i 's are positive
 - Consider all the samples (not just the misclassified ones)

MSE Margins



- If $a^T y_i = b_i$, y_i must be at distance b_i from the separating hyperplane (normalized by $\|a\|$)
- Thus b_1, b_2, \dots, b_n give relative expected distances or “*margins*” of samples from the hyperplane
- Should make b_i small if sample i is expected to be near separating hyperplane, and large otherwise
- In the absence of any additional information, set $b_1 = b_2 = \dots = b_n = 1$

MSE Matrix Notation

- Need to solve n equations
- In matrix form $\mathbf{Y}\mathbf{a}=\mathbf{b}$

$$\begin{cases} \mathbf{a}^t \mathbf{y}_1 = b_1 \\ \vdots \\ \mathbf{a}^t \mathbf{y}_n = b_n \end{cases}$$

$$\underbrace{\begin{bmatrix} \mathbf{y}_1^{(0)} & \mathbf{y}_1^{(1)} & \dots & \mathbf{y}_1^{(d)} \\ \mathbf{y}_2^{(0)} & \mathbf{y}_2^{(1)} & \dots & \mathbf{y}_2^{(d)} \\ \vdots & \vdots & & \vdots \\ \mathbf{y}_n^{(0)} & \mathbf{y}_n^{(1)} & \dots & \mathbf{y}_n^{(d)} \end{bmatrix}}_{\mathbf{Y}} \underbrace{\begin{bmatrix} \mathbf{a}_0 \\ \mathbf{a}_1 \\ \vdots \\ \mathbf{a}_d \end{bmatrix}}_{\mathbf{a}} = \underbrace{\begin{bmatrix} \mathbf{b}_1 \\ \mathbf{b}_2 \\ \vdots \\ \mathbf{b}_n \end{bmatrix}}_{\mathbf{b}}$$

Exact Solution is Rare

- Need to solve a linear system $\mathbf{Y}\mathbf{a} = \mathbf{b}$
 - \mathbf{Y} is an $n \times (d + 1)$ matrix
- Exact solution only if \mathbf{Y} is non-singular and square (the inverse \mathbf{Y}^{-1} exists)
 - $\mathbf{a} = \mathbf{Y}^{-1} \mathbf{b}$
 - (number of samples) = (number of features + 1)
 - Almost never happens in practice
 - Guaranteed to find the separating hyperplane

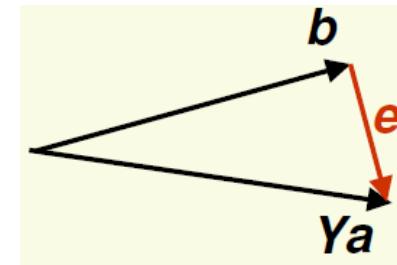
Approximate Solution

- Typically \mathbf{Y} is overdetermined, that is it has more rows (examples) than columns (features)
 - If it has more features than examples, should reduce dimensionality
- Need $\mathbf{Ya} = \mathbf{b}$, but no exact solution exists for an over-determined system of equations
 - More equations than unknowns
- Find an approximate solution
 - Note that approximate solution \mathbf{a} does **not** necessarily give the separating hyperplane in the separable case
 - But the hyperplane corresponding to \mathbf{a} may still be a good solution, especially if there is no separating hyperplane

MSE Criterion Function

- Minimum squared error approach: find \mathbf{a} which minimizes the length of the error vector \mathbf{e}

$$\mathbf{e} = \mathbf{Y}\mathbf{a} - \mathbf{b}$$



- Thus minimize the **minimum squared error criterion function**:

$$J_s(\mathbf{a}) = \|\mathbf{Y}\mathbf{a} - \mathbf{b}\|^2 = \sum_{i=1}^n (\mathbf{a}^t \mathbf{y}_i - b_i)^2$$

- Unlike the perceptron criterion function, we can optimize the minimum squared error criterion function analytically by setting the gradient to 0

Computing the Gradient

$$J_s(\mathbf{a}) = \|\mathbf{Y}\mathbf{a} - \mathbf{b}\|^2 = \sum_{i=1}^n (\mathbf{a}^t \mathbf{y}_i - b_i)^2$$

$$\begin{aligned}\nabla J_s(\mathbf{a}) &= \begin{bmatrix} \frac{\partial J_s}{\partial a_0} \\ \vdots \\ \frac{\partial J_s}{\partial a_d} \end{bmatrix} = \frac{d J_s}{d \mathbf{a}} = \sum_{i=1}^n \frac{d}{d \mathbf{a}} (\mathbf{a}^t \mathbf{y}_i - b_i)^2 \\ &= \sum_{i=1}^n 2(\mathbf{a}^t \mathbf{y}_i - b_i) \frac{d}{d \mathbf{a}} (\mathbf{a}^t \mathbf{y}_i - b_i) \\ &= \sum_{i=1}^n 2(\mathbf{a}^t \mathbf{y}_i - b_i) \mathbf{y}_i \\ &= 2 \mathbf{Y}^t (\mathbf{Y}\mathbf{a} - \mathbf{b})\end{aligned}$$

Pseudo-Inverse Solution

$$\nabla J_s(\mathbf{a}) = 2\mathbf{Y}^t(\mathbf{Y}\mathbf{a} - \mathbf{b})$$

- Setting the gradient to 0:

$$2\mathbf{Y}^t(\mathbf{Y}\mathbf{a} - \mathbf{b}) = \mathbf{0} \Rightarrow \mathbf{Y}^t\mathbf{Y}\mathbf{a} = \mathbf{Y}^t\mathbf{b}$$

- The matrix $\mathbf{Y}^t\mathbf{Y}$ is square (it has $d + 1$ rows and columns) and it is often non-singular
- If $\mathbf{Y}^t\mathbf{Y}$ is non-singular, its inverse exists and we can solve for \mathbf{a} uniquely:

$$\mathbf{a} = (\mathbf{Y}^t\mathbf{Y})^{-1}\mathbf{Y}^t\mathbf{b}$$

pseudo inverse of \mathbf{Y}
 $((\mathbf{Y}^t\mathbf{Y})^{-1}\mathbf{Y}^t)\mathbf{Y} = (\mathbf{Y}^t\mathbf{Y})^{-1}(\mathbf{Y}^t\mathbf{Y}) = I$

MSE Procedures

- Only guaranteed separating hyperplane if $\mathbf{Ya} > 0$
 - That is if all elements of vector \mathbf{Ya} are positive

$$\mathbf{Ya} = \begin{bmatrix} \mathbf{b}_1 + \boldsymbol{\varepsilon}_1 \\ \vdots \\ \mathbf{b}_n + \boldsymbol{\varepsilon}_n \end{bmatrix}$$

- where $\boldsymbol{\varepsilon}$ may be negative
- If $\boldsymbol{\varepsilon}_1, \dots, \boldsymbol{\varepsilon}_n$ are small relative to $\mathbf{b}_1, \dots, \mathbf{b}_n$, then each element of \mathbf{Ya} is positive, and \mathbf{a} gives a separating hyperplane
 - If the approximation is not good, $\boldsymbol{\varepsilon}_i$ may be large and negative, for some i , thus $\mathbf{b}_i + \boldsymbol{\varepsilon}_i$ will be negative and \mathbf{a} is not a separating hyperplane
- In linearly separable case, least squares solution \mathbf{a} does not necessarily give separating hyperplane

MSE Procedures

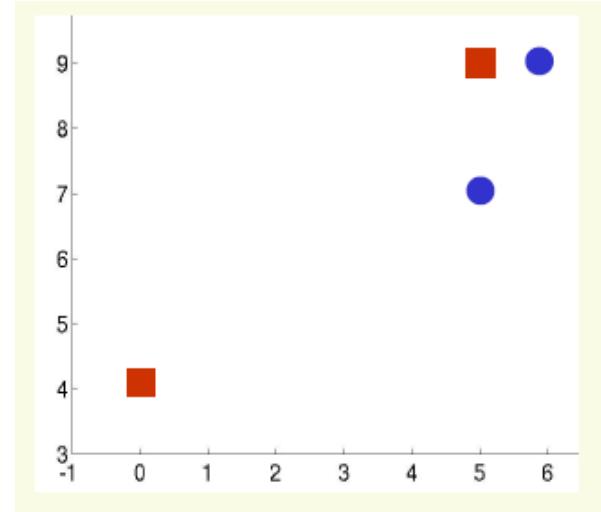
- We are free to choose \mathbf{b} . We may be tempted to make \mathbf{b} large as a way to ensure $\mathbf{Y}\mathbf{a} = \mathbf{b} > 0$
 - Does not work
 - Let β be a scalar, let's try $\beta\mathbf{b}$ instead of \mathbf{b}
- If \mathbf{a}^* is a least squares solution to $\mathbf{Y}\mathbf{a} = \mathbf{b}$, then for any scalar β , the least squares solution to $\mathbf{Y}\mathbf{a} = \beta\mathbf{b}$ is $\beta\mathbf{a}^*$

$$\arg \min_{\mathbf{a}} \|\mathbf{Y}\mathbf{a} - \beta\mathbf{b}\|^2 = \arg \min_{\mathbf{a}} \beta^2 \|\mathbf{Y}(\mathbf{a}/\beta) - \mathbf{b}\|^2 = \beta\mathbf{a}^*$$

- Thus if the i^{th} element of $\mathbf{Y}\mathbf{a}$ is less than 0, that is $\mathbf{y}_i^t \mathbf{a} < 0$, then $\mathbf{y}_i^t (\beta\mathbf{a}) < 0$,
 - The relative difference between components of \mathbf{b} matters, but not the size of each individual component

LDF using MSE: Example 1

- Class 1: (6 9), (5 7)
- Class 2: (5 9), (0 4)
- Add extra feature and “normalize”



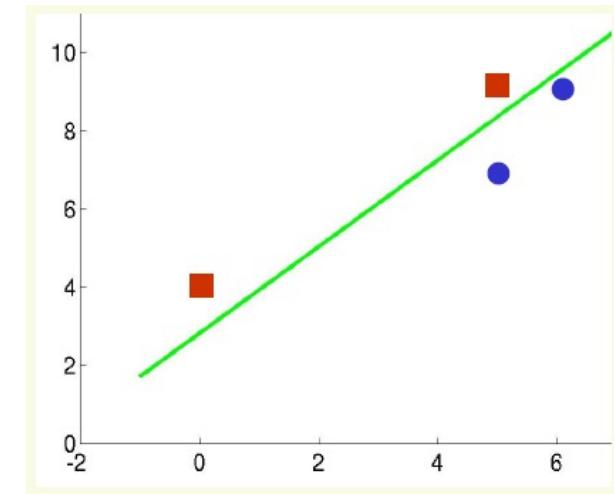
$$\mathbf{y}_1 = \begin{bmatrix} 1 \\ 6 \\ 9 \end{bmatrix} \quad \mathbf{y}_2 = \begin{bmatrix} 1 \\ 5 \\ 7 \end{bmatrix} \quad \mathbf{y}_3 = \begin{bmatrix} -1 \\ -5 \\ -9 \end{bmatrix} \quad \mathbf{y}_4 = \begin{bmatrix} -1 \\ 0 \\ -4 \end{bmatrix}$$

$$\mathbf{Y} = \begin{bmatrix} 1 & 6 & 9 \\ 1 & 5 & 7 \\ -1 & -5 & -9 \\ -1 & 0 & -4 \end{bmatrix}$$

LDF using MSE: Example 1

- Choose $\mathbf{b} = [1 \ 1 \ 1 \ 1]^T$
- In Matlab, $\mathbf{a} = \mathbf{Y} \setminus \mathbf{b}$ solves the least squares problem

$$\mathbf{a} = \begin{bmatrix} 2.66 \\ 1.045 \\ -0.944 \end{bmatrix}$$

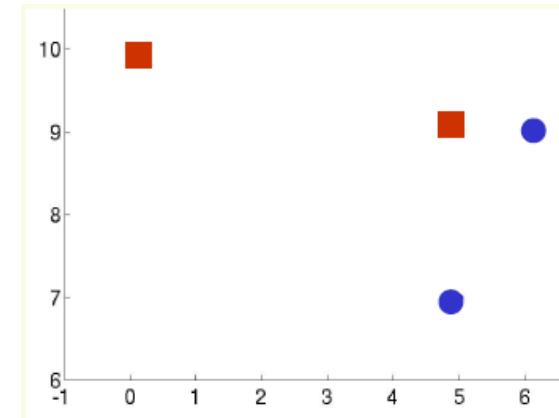


- Note \mathbf{a} is an approximation to $\mathbf{Ya} = \mathbf{b}$, since no exact solution exists
- This solution gives a separating hyperplane since $\mathbf{Ya} > 0$

$$\mathbf{Ya} = \begin{bmatrix} 0.44 \\ 1.28 \\ 0.61 \\ 1.11 \end{bmatrix}$$

LDF using MSE: Example 2

- Class 1: (6 9), (5 7)
- Class 2: (5 9), (0 10)
- The last sample is very far compared to others from the separating hyperplane

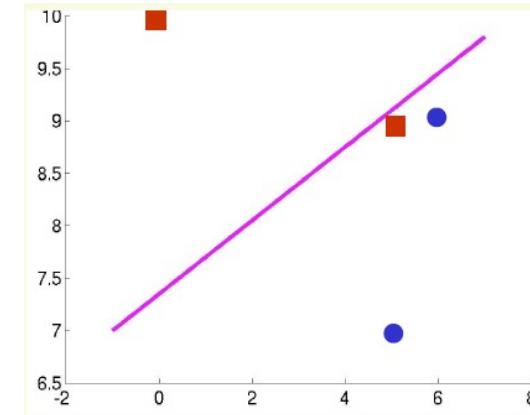


$$\mathbf{y}_1 = \begin{bmatrix} 1 \\ 6 \\ 9 \end{bmatrix} \quad \mathbf{y}_2 = \begin{bmatrix} 1 \\ 5 \\ 7 \end{bmatrix} \quad \mathbf{y}_3 = \begin{bmatrix} -1 \\ -5 \\ -9 \end{bmatrix} \quad \mathbf{y}_4 = \begin{bmatrix} -1 \\ 0 \\ -10 \end{bmatrix}$$

$$\mathbf{Y} = \begin{bmatrix} 1 & 6 & 9 \\ 1 & 5 & 7 \\ -1 & -5 & -9 \\ -1 & 0 & -10 \end{bmatrix}$$

LDF using MSE: Example 2

- Choose $\mathbf{b} = [1 \ 1 \ 1 \ 1]^T$
- In Matlab, $\mathbf{a} = \mathbf{Y} \setminus \mathbf{b}$ solves the least squares problem



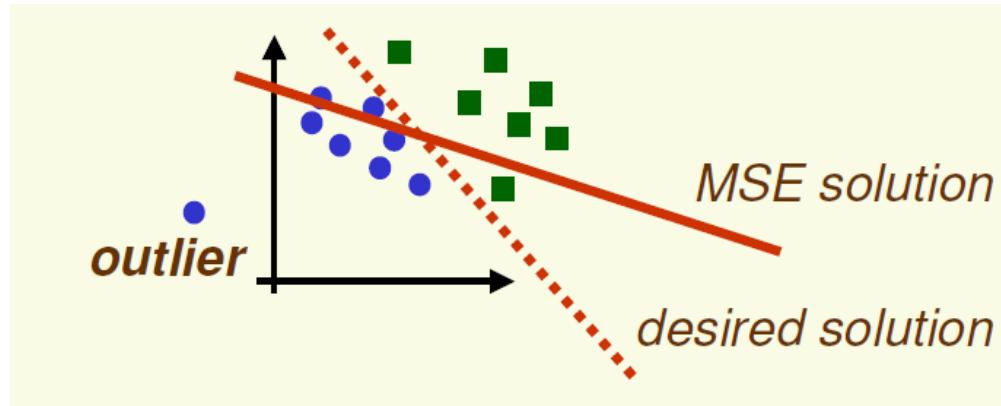
$$\mathbf{a} = \begin{bmatrix} 3.2 \\ 0.2 \\ -0.4 \end{bmatrix}$$

$$\mathbf{Y}\mathbf{a} = \begin{bmatrix} 0.2 \\ 0.9 \\ -0.04 \\ 1.16 \end{bmatrix} \neq \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

- This solution does not provide a separating hyperplane since $\mathbf{a}^T \mathbf{y}_3 < 0$

LDF using MSE: Example 2

- MSE pays too much attention to isolated “noisy” examples
 - such examples are called outliers

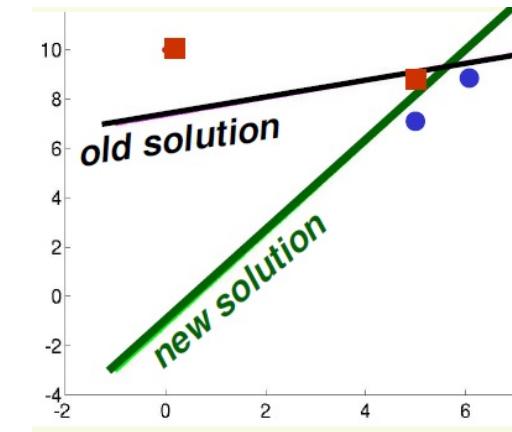


- No problems with convergence
- Solution ranges from reasonable to good

LDF using MSE: Example 2

- We can see that the 4th point is vary far from separating hyperplane
 - In practice we don't know this
- A more appropriate \mathbf{b} could be
- In Matlab, solve $\mathbf{a} = \mathbf{Y} \setminus \mathbf{b}$

$$\mathbf{b} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 10 \end{bmatrix}$$



$$\mathbf{a} = \begin{bmatrix} -1.1 \\ 1.7 \\ -0.9 \end{bmatrix}$$

$$\mathbf{Y}\mathbf{a} = \begin{bmatrix} 0.9 \\ 1.0 \\ 0.8 \\ 10.0 \end{bmatrix} \neq \begin{bmatrix} 1 \\ 1 \\ 1 \\ 10 \end{bmatrix}$$

- This solution gives the separating hyperplane since $\mathbf{Y}\mathbf{a} > 0$

Gradient Descent for MSE

$$J_s(a) = \|Ya - b\|^2$$

- May wish to find MSE solution by gradient descent:
 1. Computing the inverse of $Y^t Y$ may be too costly
 2. $Y^t Y$ may be close to singular if samples are highly correlated (rows of Y are almost linear combinations of each other) computing the inverse of $Y^t Y$ is not numerically stable
- As shown before, the gradient is:

$$\nabla J_s(a) = 2Y^t(Ya - b)$$

Widrow-Hoff Procedure

$$\nabla J_s(\mathbf{a}) = 2\mathbf{Y}^t(\mathbf{Y}\mathbf{a} - \mathbf{b})$$

- Thus the update rule for gradient descent is:

$$\mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} - \eta^{(k)} \mathbf{Y}^t (\mathbf{Y}\mathbf{a}^{(k)} - \mathbf{b})$$

- If $\eta^{(k)}=\eta^{(1)}/k$, then $\mathbf{a}^{(k)}$ converges to the MSE solution \mathbf{a} , that is $\mathbf{Y}^t(\mathbf{Y}\mathbf{a}-\mathbf{b})=0$
- The *Widrow-Hoff procedure* reduces storage requirements by considering single samples sequentially

$$\mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} - \eta^{(k)} \mathbf{y}_i (\mathbf{y}_i^t \mathbf{a}^{(k)} - b_i)$$

LDF Summary

- **Perceptron procedures**
 - Find a separating hyperplane in the linearly separable case,
 - Do not converge in the non-separable case
 - Can force convergence by using a decreasing learning rate, but are not guaranteed a reasonable stopping point
- **MSE procedures**
 - Converge in separable and not separable case
 - May not find separating hyperplane even if classes are linearly separable
 - Use pseudoinverse if $\mathbf{Y}^t\mathbf{Y}$ is not singular and not too large
 - Use gradient descent (Widrow-Hoff procedure) otherwise

Support Vector Machines

SVM Resources

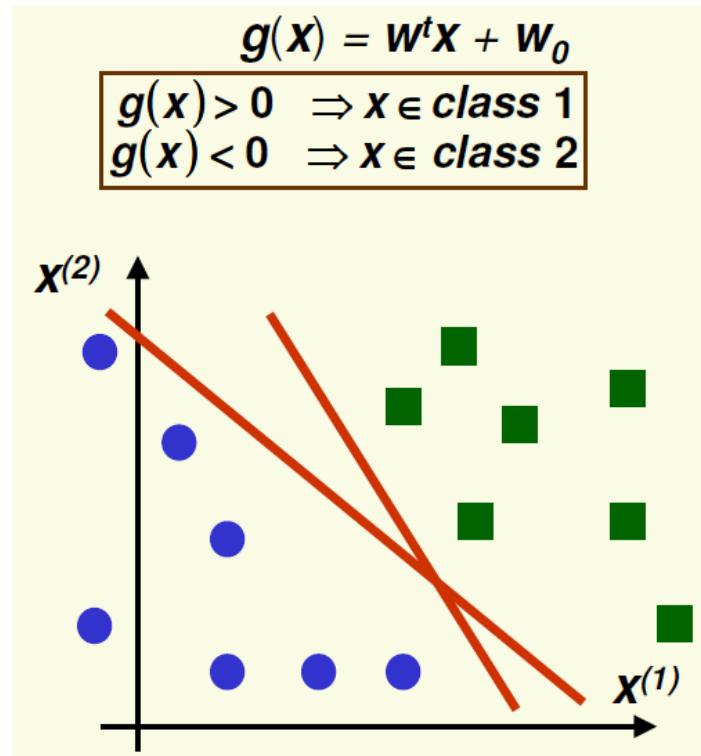
- Burges tutorial
 - <http://research.microsoft.com/en-us/um/people/cburges/papers/SVMTutorial.pdf>
- Shawe-Taylor and Christianini tutorial
 - <http://www.support-vector.net/icml-tutorial.pdf>
- Lib SVM
 - <http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
- LibLinear
 - <http://www.csie.ntu.edu.tw/~cjlin/liblinear/>
- SVM Light
 - <http://svmlight.joachims.org/>
- Power Mean SVM (very fast for histogram features)
 - <https://sites.google.com/site/wujx2001/home/power-mean-svm>

SVMs

- One of the most important developments in pattern recognition in the last decades
- Elegant theory
 - Has good generalization properties
- Have been applied to diverse problems very successfully

Linear Discriminant Functions

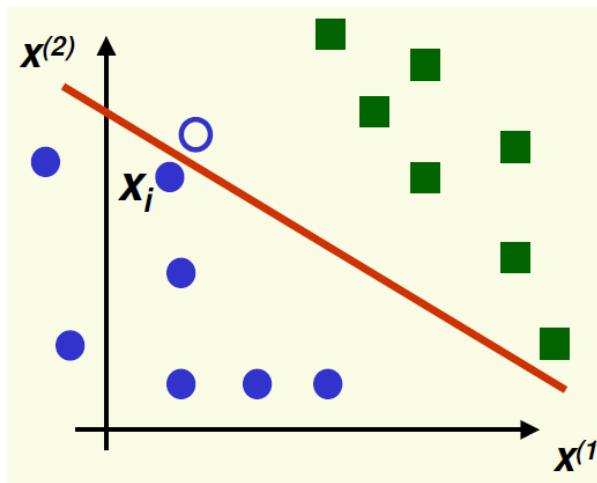
- A discriminant function is linear if it can be written as



- which separating hyperplane should we choose?

Linear Discriminant Functions

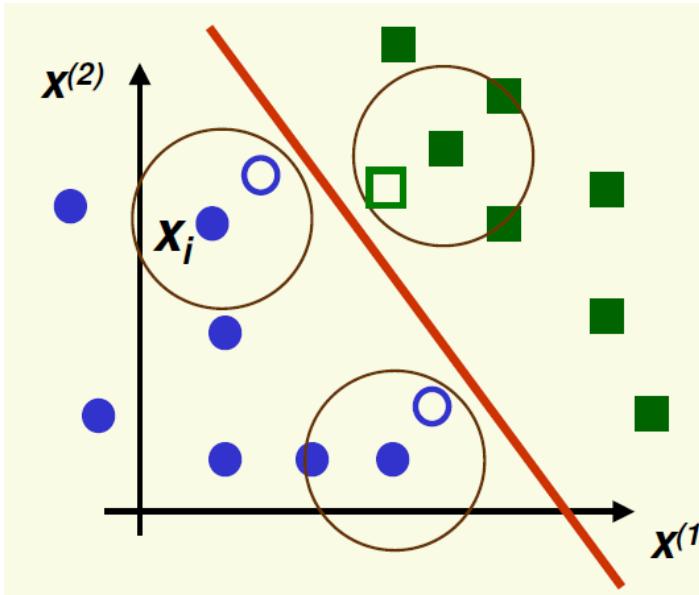
- Training data is just a subset of all possible data
 - Suppose hyperplane is close to sample \mathbf{x}_i
 - If we see new sample close to \mathbf{x}_i , it may be on the wrong side of the hyperplane



- Poor generalization (performance on unseen data)

Linear Discriminant Functions

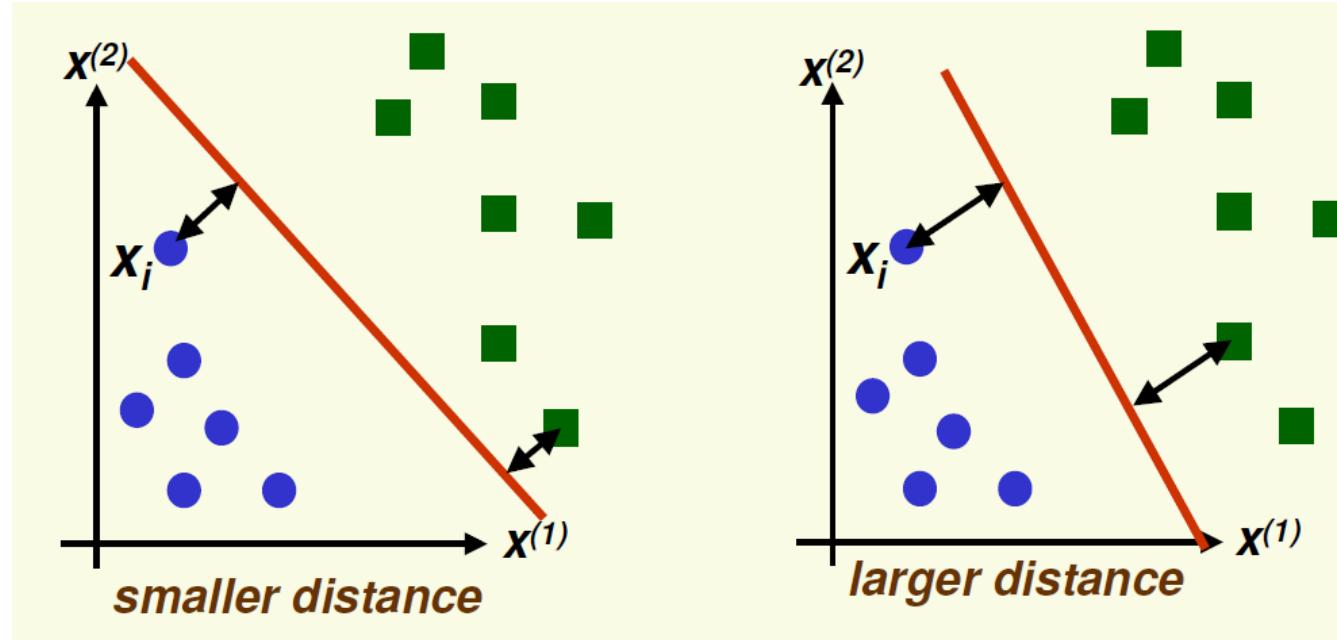
- Hyperplane as far as possible from any sample



- New samples close to the old samples will be classified correctly
- Good generalization

SVM

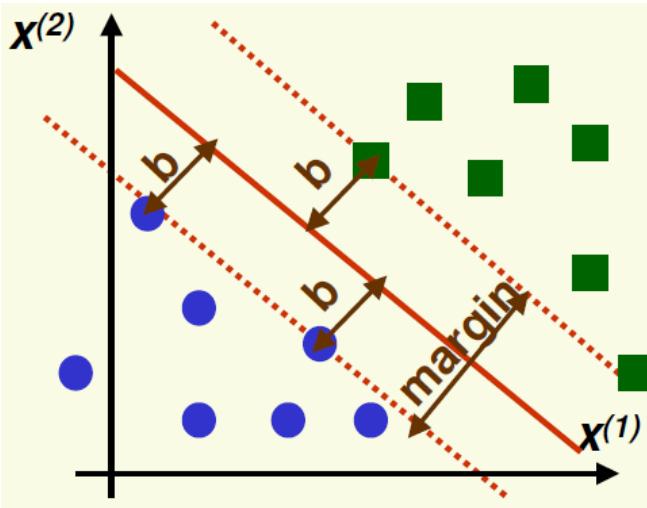
- Idea: maximize distance to the closest example



- For the optimal hyperplane
 - distance to the closest negative example = distance to the closest positive example

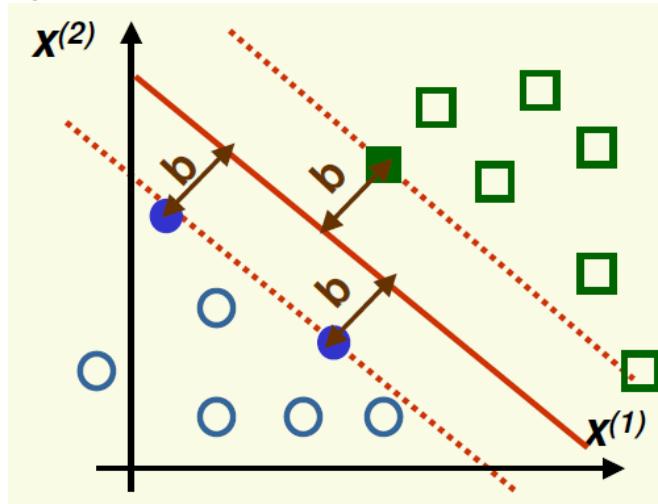
SVM: Linearly Separable Case

- SVM: maximize the margin



- The *margin* is twice the absolute value of distance b of the closest example to the separating hyperplane
- Better generalization (performance on test data)
 - in practice
 - and in theory

SVM: Linearly Separable Case



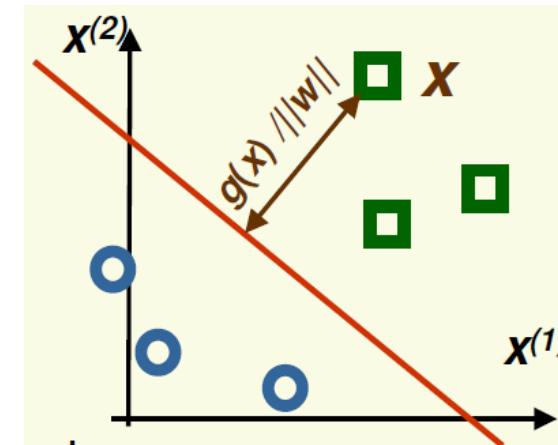
- **Support vectors** are the samples closest to the separating hyperplane
 - They are the most difficult patterns to classify
 - Recall perceptron update rule
- Optimal hyperplane is completely defined by support vectors
 - Of course, we do not know which samples are support vectors without finding the optimal hyperplane

SVM: Formula for the Margin

$$g(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

- Absolute distance between \mathbf{x} and the boundary $g(\mathbf{x}) = 0$

$$\frac{|\mathbf{w}^t \mathbf{x} + w_0|}{\|\mathbf{w}\|}$$



- Distance is unchanged for hyperplane

$$g_1(\mathbf{x}) = \alpha g(\mathbf{x}) \quad \frac{|\alpha \mathbf{w}^t \mathbf{x} + \alpha w_0|}{\|\alpha \mathbf{w}\|} = \frac{|\mathbf{w}^t \mathbf{x} + w_0|}{\|\mathbf{w}\|}$$

- Let \mathbf{x}_i be an example closest to the boundary (on the positive side). Set:

$$|\mathbf{w}^t \mathbf{x}_i + w_0| = 1$$

- Now the largest margin hyperplane is unique

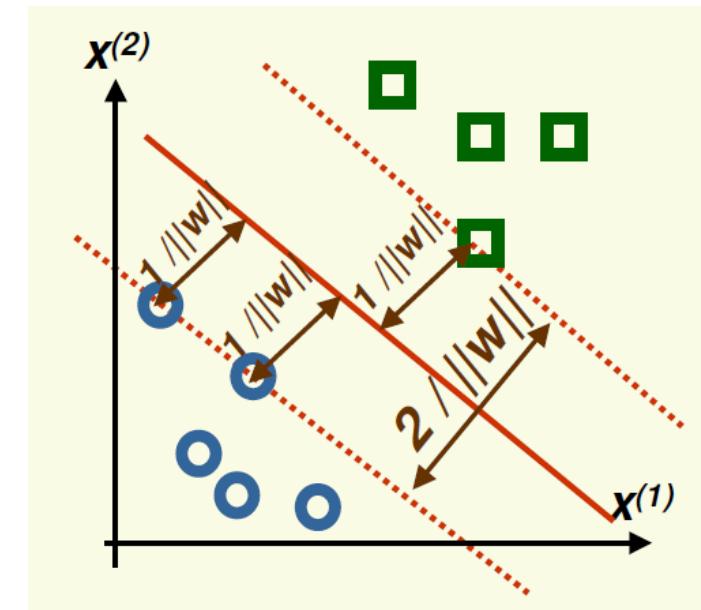
SVM: Formula for the Margin

- For uniqueness, set $|\mathbf{w}^T \mathbf{x}_i + w_0| = 1$ for any sample \mathbf{x}_i closest to the boundary
- The distance from closest sample \mathbf{x}_i to $g(\mathbf{x}) = 0$ is

$$\frac{|\mathbf{w}^T \mathbf{x}_i + w_0|}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}$$

- Thus the margin is

$$m = \frac{2}{\|\mathbf{w}\|}$$



SVM: Optimal Hyperplane

- Maximize margin

$$m = \frac{2}{\|w\|}$$

- Subject to constraints

$$\begin{cases} w^t x_i + w_0 \geq 1 & \text{if } x_i \text{ is positive example} \\ w^t x_i + w_0 \leq -1 & \text{if } x_i \text{ is negative example} \end{cases}$$

- Let

$$\begin{cases} z_i = 1 & \text{if } x_i \text{ is positive example} \\ z_i = -1 & \text{if } x_i \text{ is negative example} \end{cases}$$

- Can convert our problem to minimize

$$\text{minimize } J(w) = \frac{1}{2} \|w\|^2$$

$$\text{constrained to } z_i (w^t x_i + w_0) \geq 1 \quad \forall i$$

- $J(w)$ is a quadratic function, thus there is a single global minimum

SVM: Optimal Hyperplane

- Use Kuhn-Tucker theorem to convert our problem to:
 - Also known as the Karush–Kuhn–Tucker theorem, i.e., the KKT theorem

$$\begin{aligned} & \text{maximize} && L_D(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j z_i z_j x_i^t x_j \\ & \text{constrained to} && \alpha_i \geq 0 \quad \forall i \quad \text{and} \quad \sum_{i=1}^n \alpha_i z_i = 0 \end{aligned}$$

- $\alpha = \{\alpha_1, \dots, \alpha_n\}$ are new variables, one for each sample
- Optimized by quadratic programming

SVM: Optimal Hyperplane

- After finding the optimal $\alpha = \{\alpha_1, \dots, \alpha_n\}$
- Final discriminant function:

$$g(\mathbf{x}) = \left(\sum_{x_i \in S} \alpha_i z_i \mathbf{x}_i \right)^t \mathbf{x} + w_0$$

- where S is the set of support vectors

$$S = \{x_i \mid \alpha_i \neq 0\}$$

SVM: Optimal Hyperplane

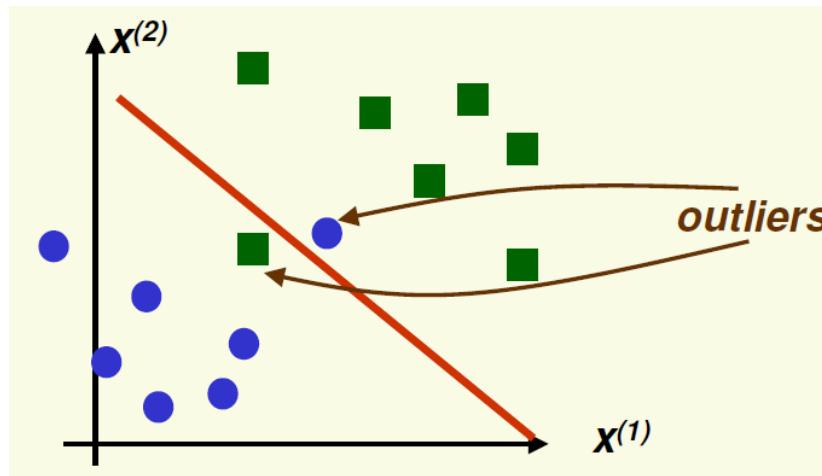
$$\text{maximize } L_D(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j z_i z_j x_i^t x_j$$

$$\text{constrained to } \alpha_i \geq 0 \quad \forall i \quad \text{and} \quad \sum_{i=1}^n \alpha_i z_i = 0$$

- $L_D(\alpha)$ depends on the number of samples, not on dimension
 - samples appear only through the dot products $x_i^t x_i$
 - This will become important when looking for a nonlinear discriminant function, as we will see soon

SVM: Non-Separable Case

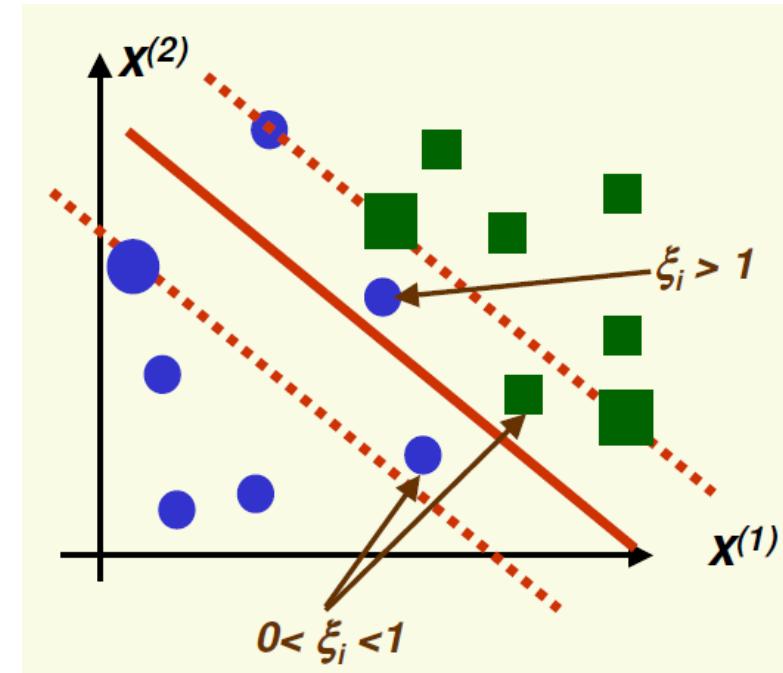
- Data are most likely to be not linearly separable, but linear classifier may still be appropriate



- Can apply SVM in non linearly separable case
- Data should be “almost” linearly separable for good performance

SVM: Non-Separable Case

- Use **slack variables** ξ_1, \dots, ξ_n (one for each sample)
- Change constraints from $z_i(w^t x_i + w_o) \geq 1 \quad \forall i$ to
$$z_i(w^t x_i + w_o) \geq 1 - \xi_i \quad \forall i$$
- ξ_i is a measure of deviation from the ideal for x_i
 - $\xi_i > 1$: x_i is on the wrong side of the separating hyperplane
 - $0 < \xi_i < 1$: x_i is on the right side of separating hyperplane but within the region of maximum margin
 - $\xi_i < 0$: is the ideal case for x_i



SVM: Non-Separable Case

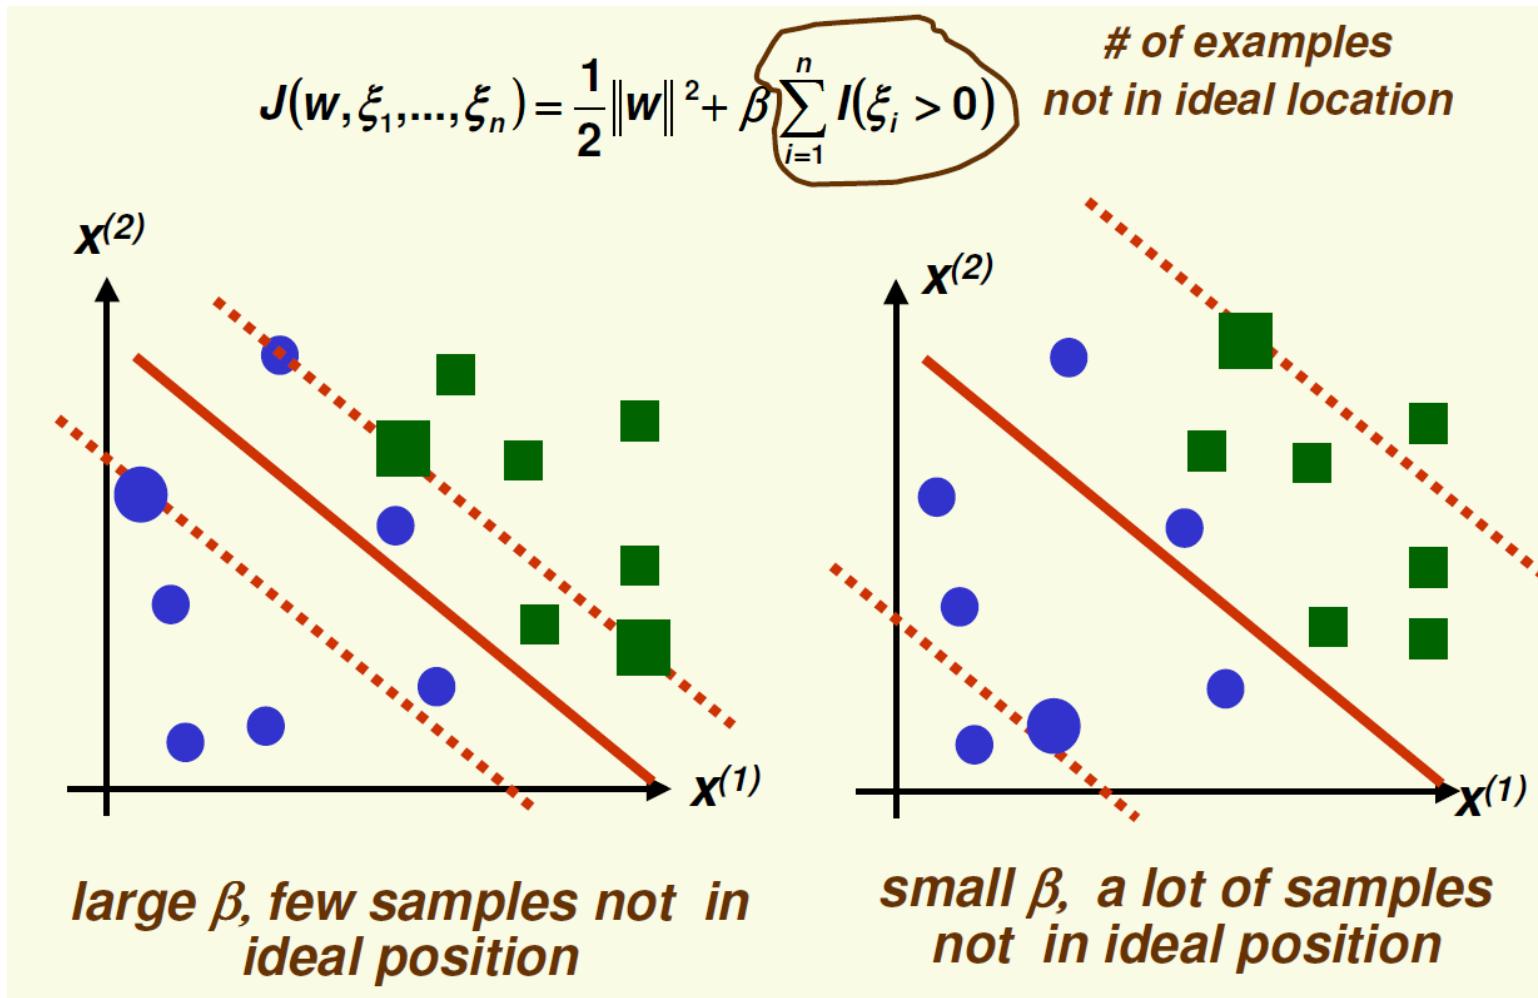
- We would like to minimize

$$J(\mathbf{w}, \xi_1, \dots, \xi_n) = \frac{1}{2} \|\mathbf{w}\|^2 + \beta \sum_{i=1}^n I(\xi_i > 0)$$

*# of samples
not in ideal location*

- where $I(\xi_i > 0) = \begin{cases} 1 & \text{if } \xi_i > 0 \\ 0 & \text{if } \xi_i \leq 0 \end{cases}$
- Constrained to $\mathbf{z}_i(\mathbf{w}^t \mathbf{x}_i + w_0) \geq 1 - \xi_i$ and $\xi_i \geq 0 \quad \forall i$
- β is a constant that measures the relative weight of first and second term
 - If β is small, we allow a lot of samples to be in not ideal positions
 - If β is large, few samples can be in non-ideal positions

SVM: Non-Separable Case



SVM: Non-Separable Case

- Unfortunately this minimization problem is NP-hard due to the discontinuity of $I(\xi_i)$
- Instead, we minimize

$$J(w, \xi_1, \dots, \xi_n) = \frac{1}{2} \|w\|^2 + \beta \sum_{i=1}^n \xi_i$$

*a measure of
of misclassified
examples*

- Subject to

$$\begin{cases} z_i (w^t x_i + w_0) \geq 1 - \xi_i & \forall i \\ \xi_i \geq 0 & \forall i \end{cases}$$

SVM: Non-Separable Case

- Use Kuhn-Tucker theorem to convert to:

$$\begin{aligned} \text{maximize } L_D(\alpha) &= \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j \mathbf{z}_i \mathbf{z}_j \mathbf{x}_i^t \mathbf{x}_j \\ \text{constrained to } & 0 \leq \alpha_i \leq \beta \quad \forall i \quad \text{and} \quad \sum_{i=1}^n \alpha_i \mathbf{z}_i = 0 \end{aligned}$$

- \mathbf{w} is computed using:

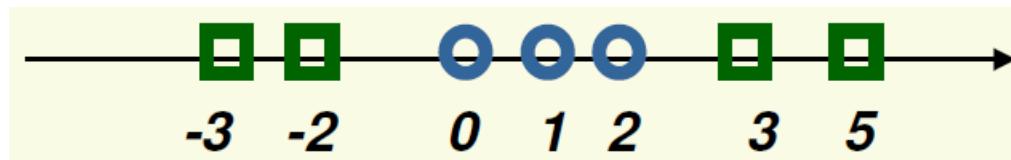
$$\mathbf{w} = \sum_{i=1}^n \alpha_i \mathbf{z}_i \mathbf{x}_i$$

- Remember that

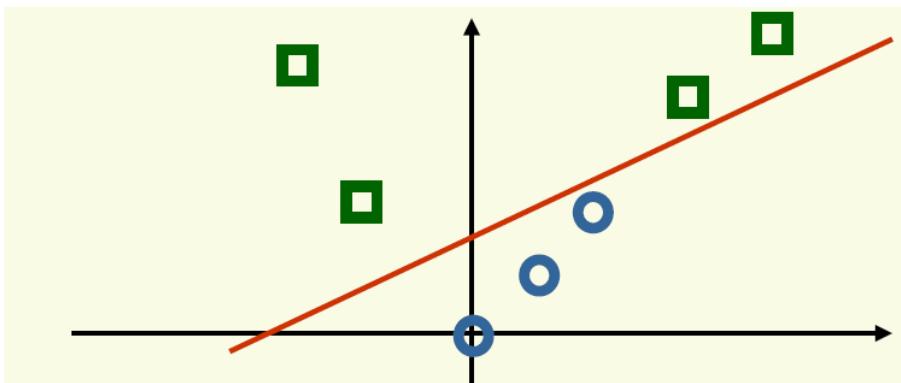
$$g(\mathbf{x}) = \left(\sum_{\mathbf{x}_i \in S} \alpha_i \mathbf{z}_i \mathbf{x}_i \right)^t \mathbf{x} + w_0$$

Nonlinear Mapping

- Cover's theorem: “*a pattern-classification problem cast in a high dimensional space non-linearly is more likely to be linearly separable than in a low-dimensional space*”
- One dimensional space, not linearly separable

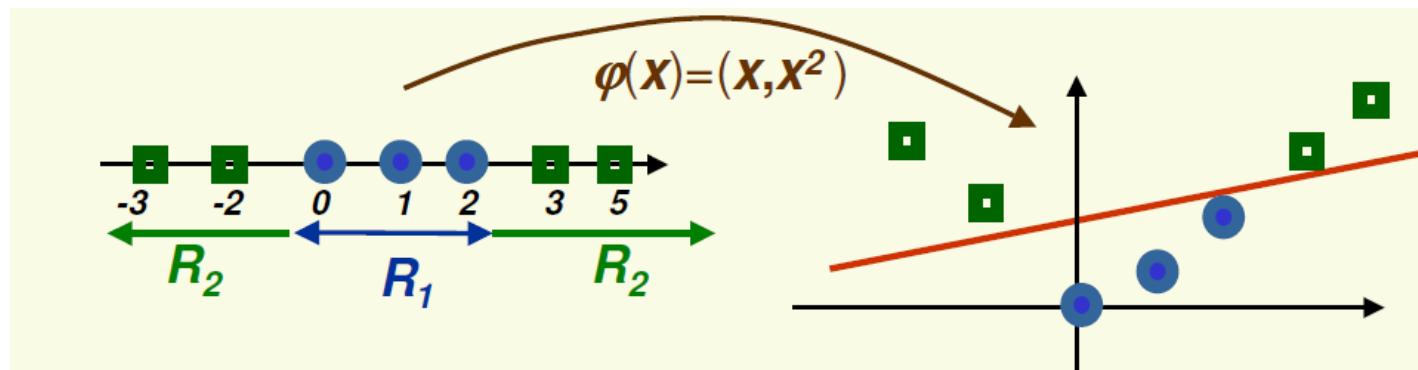


- Lift to two dimensional space with $\phi(x)=(x,x^2)$



Nonlinear Mapping

- To solve a non linear classification problem with a linear classifier
- 1. Project data x to high dimension using function $\phi(x)$
- 2. Find a linear discriminant function for transformed data $\phi(x)$
- 3. Final nonlinear discriminant function is $g(x) = w^t \phi(x) + w_0$



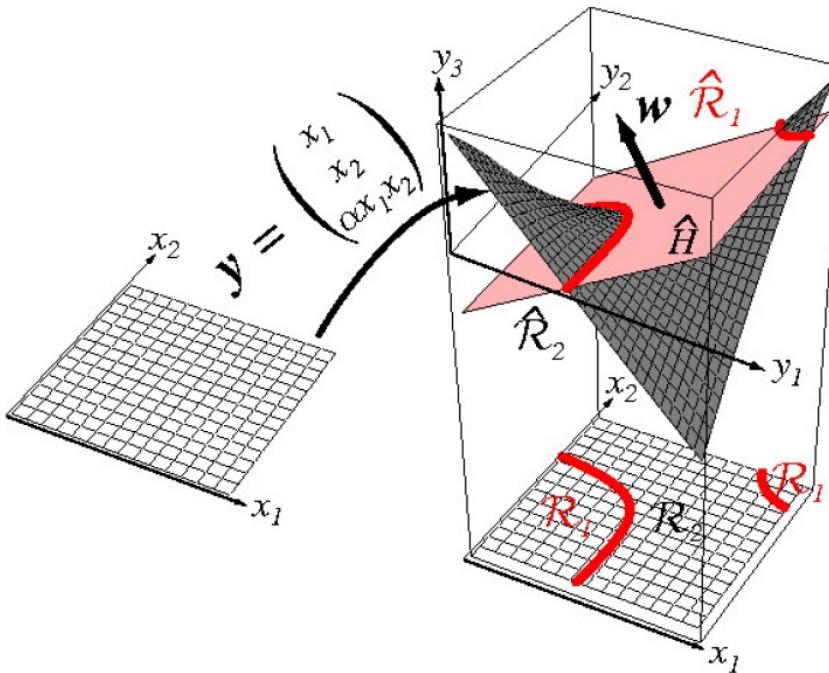
- In 2D, the discriminant function is linear

$$g\left(\begin{bmatrix} x^{(1)} \\ x^{(2)} \end{bmatrix}\right) = [w_1 \quad w_2] \begin{bmatrix} x^{(1)} \\ x^{(2)} \end{bmatrix} + w_0$$

- In 1D, the discriminant function is not linear

$$g(x) = w_1 x + w_2 x^2 + w_0$$

Nonlinear Mapping



- However, there always exists a mapping of N samples to an N-dimensional space in which the samples are separable by hyperplanes

Nonlinear SVM

- Can use any linear classifier after lifting data to a higher dimensional space. However we will have to deal with the curse of dimensionality
 - Poor generalization to test data
 - Computationally expensive
- SVM avoids the curse of dimensionality problems
 - Enforcing largest margin permits good generalization
 - It can be shown that generalization in SVM is a function of the margin, independent of the dimensionality
 - Computation in the higher dimensional case is performed only implicitly through the use of *kernel functions*

Kernels

- SVM optimization:

maximize

$$L_D(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j z_i z_j \mathbf{x}_i^t \mathbf{x}_j$$

- Note this optimization depends on samples \mathbf{x}_i only through the dot product $\mathbf{x}_i^t \mathbf{x}_j$
- If we lift \mathbf{x}_i to high dimension using $\varphi(\mathbf{x})$, we need to compute high dimensional product $\varphi(\mathbf{x}_i)^t \varphi(\mathbf{x}_j)$

maximize

$$L_D(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j z_i z_j \varphi(\mathbf{x}_i)^t \varphi(\mathbf{x}_j)$$

$K(\mathbf{x}_i, \mathbf{x}_j)$

- Idea: find kernel function $K(\mathbf{x}_i, \mathbf{x}_j)$ s.t. $K(\mathbf{x}_i, \mathbf{x}_j) = \varphi(\mathbf{x}_i)^t \varphi(\mathbf{x}_j)$

Kernel Trick

- Then we only need to compute $K(x_i, x_j)$ instead of $\phi(x_i)^t \phi(x_j)$
- “**kernel trick**”: do not need to perform operations in high dimensional space explicitly

Kernel Example

- Suppose we have two features and $K(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^t \mathbf{y})^2$
- Which mapping $\Phi(\mathbf{x})$ does this correspond to?

$$\begin{aligned} K(\mathbf{x}, \mathbf{y}) &= (\mathbf{x}^t \mathbf{y})^2 = \left(\begin{bmatrix} \mathbf{x}^{(1)} & \mathbf{x}^{(2)} \end{bmatrix} \begin{bmatrix} \mathbf{y}^{(1)} \\ \mathbf{y}^{(2)} \end{bmatrix} \right)^2 = (\mathbf{x}^{(1)} \mathbf{y}^{(1)} + \mathbf{x}^{(2)} \mathbf{y}^{(2)})^2 \\ &= (\mathbf{x}^{(1)} \mathbf{y}^{(1)})^2 + 2(\mathbf{x}^{(1)} \mathbf{y}^{(1)})(\mathbf{x}^{(2)} \mathbf{y}^{(2)}) + (\mathbf{x}^{(2)} \mathbf{y}^{(2)})^2 \\ &= \begin{bmatrix} (\mathbf{x}^{(1)})^2 & \sqrt{2} \mathbf{x}^{(1)} \mathbf{x}^{(2)} & (\mathbf{x}^{(2)})^2 \end{bmatrix} \begin{bmatrix} (\mathbf{y}^{(1)})^2 & \sqrt{2} \mathbf{y}^{(1)} \mathbf{y}^{(2)} & (\mathbf{y}^{(2)})^2 \end{bmatrix}^t \end{aligned}$$

$$\varphi(\mathbf{x}) = \begin{bmatrix} (\mathbf{x}^{(1)})^2 & \sqrt{2} \mathbf{x}^{(1)} \mathbf{x}^{(2)} & (\mathbf{x}^{(2)})^2 \end{bmatrix}$$

Choice of Kernel

- How to choose kernel function $K(x_i, x_j)$?
 - $K(x_i, x_j)$ should correspond to $\phi(x_i)^t \phi(x_j)$ in a higher dimensional space
 - Mercer's condition tells us which kernel function can be expressed as dot product of two vectors
 - If K and K' are kernels $aK+bK'$ is a kernel
- Intuitively: Kernel should measure the similarity between x_i and x_j
 - As inner product measures similarity of unit vectors
 - May be problem-specific

Choice of Kernel

- Some common choices:
 - Polynomial kernel

$$K(x_i, x_j) = (x_i^t x_j + 1)^p$$

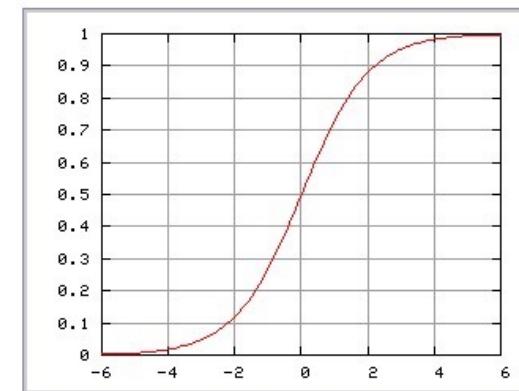
- Gaussian radial Basis kernel

$$K(x_i, x_j) = \exp\left(-\frac{1}{2\sigma^2} \|x_i - x_j\|^2\right)$$

- Hyperbolic tangent (sigmoid) kernel

$$K(x_i, x_j) = \tanh(k x_i^t x_j + c)$$

- The mappings $\Phi(x_i)$ never have to be computed!!



Intersection Kernel

- Feature vectors are histograms

$$K(x_i, x_j) = \sum_{k=1}^n \min(x_{ik}, x_{jk})$$

- When $K(x_i, x_j)$ is small, x_i and x_j are dissimilar
- When $K(x_i, x_j)$ is large, x_i and x_j are similar
- The mapping $\Phi(x)$ does not exist

More Additive Kernels

- χ^2 kernel

$$K_{\chi^2} = \sum_{k=1}^n \frac{2x_k y_k}{x_k + y_k}$$

- Hellinger's kernel

$$K_H = \sum_{k=1}^n \sqrt{x_k y_k}$$

- Designed for feature vectors that are histograms
 - Can be used for other feature vectors
- Offer very large speed-ups

The Kernel Matrix

- a.k.a the Gram matrix

K =

K(1,1)	K(1,2)	K(1,3)	...	K(1,m)
K(2,1)	K(2,2)	K(2,3)	...	K(2,m)
...
K(m,1)	K(m,2)	K(m,3)	...	K(m,m)

- Contains all necessary information for the learning algorithm
- Fuses information about the data and the kernel (similarity measure)

Bad Kernels

- The kernel matrix is mostly diagonal
 - All points are orthogonal to each other
- Bad similarity measure
- Too many irrelevant features in high dimensional space
- We need problem-specific knowledge to choose appropriate kernel

Nonlinear SVM Step-by-Step

- Start with data $\mathbf{x}_1, \dots, \mathbf{x}_n$ which live in feature space of dimension d
- Choose kernel $K(\mathbf{x}_i, \mathbf{x}_j)$ or function $\Phi(\mathbf{x}_i)$ which lifts sample \mathbf{x}_i to a higher dimensional space
- Find the maximum margin linear discriminant function in the higher dimensional space by using quadratic programming package to solve:

$$\text{maximize } L_D(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j z_i z_j K(\mathbf{x}_i, \mathbf{x}_j)$$

constrained to $0 \leq \alpha_i \leq \beta \quad \forall i$ and $\sum_{i=1}^n \alpha_i z_i = 0$

Nonlinear SVM Step-by-Step

- Weight vector w in the high dimensional space:

$$w = \sum_{x_i \in S} \alpha_i z_i \varphi(x_i)$$

- where S is the set of support vectors
- Linear discriminant function of maximum margin in the high dimensional space:

$$g(\varphi(x)) = w^t \varphi(x) = \left(\sum_{x_i \in S} \alpha_i z_i \varphi(x_i) \right)^t \varphi(x)$$

- Non linear discriminant function in the original space:

$$g(x) = \left(\sum_{x_i \in S} \alpha_i z_i \varphi(x_i) \right)^t \varphi(x) = \sum_{x_i \in S} \alpha_i z_i \varphi^t(x_i) \varphi(x) = \sum_{x_i \in S} \alpha_i z_i K(x_i, x)$$

- decide class 1 if $g(x) > 0$, otherwise decide class 2

Nonlinear SVM

- Nonlinear discriminant function

$$g(\mathbf{x}) = \sum_{x_i \in S} [\alpha_i] [\mathbf{z}_i] [K(x_i, \mathbf{x})]$$

$$g(\mathbf{x}) = \sum$$

weight of support
vector \mathbf{x}_i

∓ 1

“inverse distance”
from \mathbf{x} to
support vector \mathbf{x}_i

most important
training samples,
i.e. support vectors

$$K(x_i, \mathbf{x}) = \exp\left(-\frac{1}{2\sigma^2} \|\mathbf{x}_i - \mathbf{x}\|^2\right)$$

SVM Example: XOR Problem

- Class 1: $\mathbf{x}_1 = [1, -1]$, $\mathbf{x}_2 = [-1, 1]$
- Class 2: $\mathbf{x}_3 = [1, 1]$, $\mathbf{x}_4 = [-1, -1]$
- Use polynomial kernel of degree 2:

$$K(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i^t \mathbf{x}_j + 1)^2$$

- This kernel corresponds to the mapping

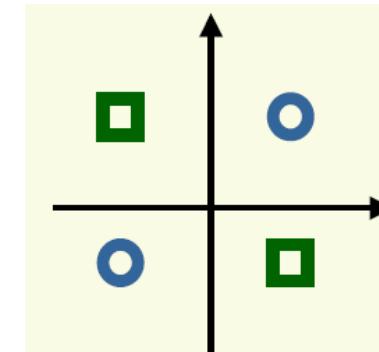
$$\varphi(\mathbf{x}) = [1 \quad \sqrt{2}\mathbf{x}^{(1)} \quad \sqrt{2}\mathbf{x}^{(2)} \quad \sqrt{2}\mathbf{x}^{(1)}\mathbf{x}^{(2)} \quad (\mathbf{x}^{(1)})^2 \quad (\mathbf{x}^{(2)})^2]^t$$

- Need to maximize

$$L_D(\alpha) = \sum_{i=1}^4 \alpha_i - \frac{1}{2} \sum_{i=1}^4 \sum_{j=1}^4 \alpha_i \alpha_j \mathbf{z}_i \mathbf{z}_j (\mathbf{x}_i^t \mathbf{x}_j + 1)^2$$

constrained to

$$0 \leq \alpha_i \quad \forall i \quad \text{and} \quad \alpha_1 + \alpha_2 - \alpha_3 - \alpha_4 = 0$$



SVM Example: XOR Problem

- After some manipulation ...
- The solution is $\alpha_1 = \alpha_2 = \alpha_3 = \alpha_4 = 0.25$
 - satisfies the constraints

$$\forall i, 0 \leq \alpha_i \text{ and } \alpha_1 + \alpha_2 - \alpha_3 - \alpha_4 = 0$$

- All samples are support vectors

SVM Example: XOR Problem

$$\varphi(\mathbf{x}) = [1 \ \sqrt{2}\mathbf{x}^{(1)} \ \sqrt{2}\mathbf{x}^{(2)} \ \sqrt{2}\mathbf{x}^{(1)}\mathbf{x}^{(2)} \ (\mathbf{x}^{(1)})^2 \ (\mathbf{x}^{(2)})^2]^t$$

- The weight vector \mathbf{w} is:

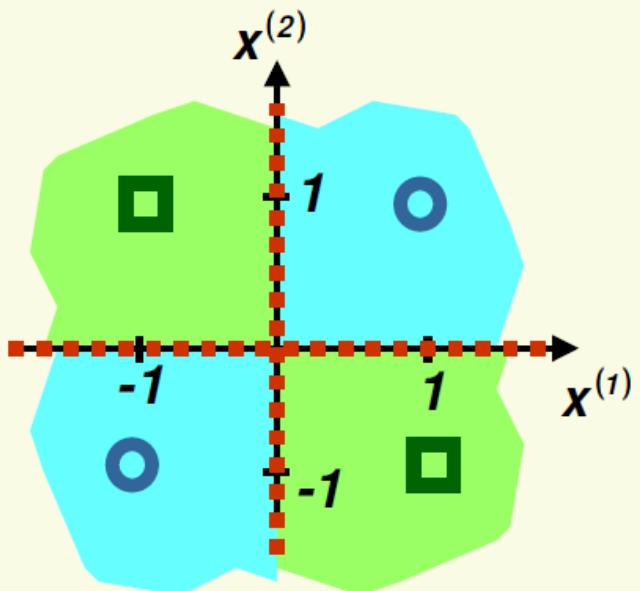
$$\begin{aligned}\mathbf{w} &= \sum_{i=1}^4 \alpha_i z_i \varphi(\mathbf{x}_i) = 0.25(\varphi(\mathbf{x}_1) + \varphi(\mathbf{x}_2) - \varphi(\mathbf{x}_3) - \varphi(\mathbf{x}_4)) \\ &= [0 \ 0 \ 0 \ -\sqrt{2} \ 0 \ 0]\end{aligned}$$

- Thus the nonlinear discriminant function is:

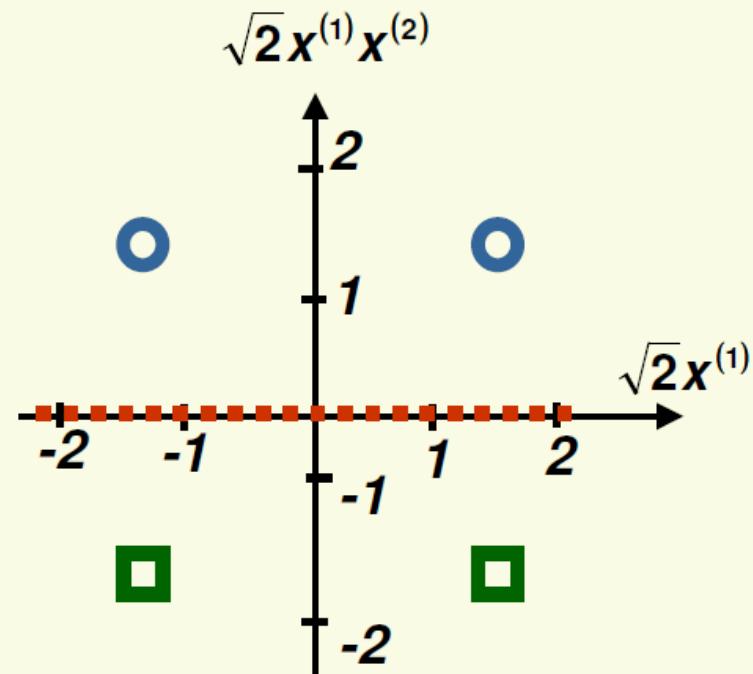
$$g(\mathbf{x}) = \mathbf{w} \varphi(\mathbf{x}) = \sum_{i=1}^6 w_i \varphi_i(\mathbf{x}) = -\sqrt{2}(\sqrt{2}\mathbf{x}^{(1)}\mathbf{x}^{(2)}) = -2\mathbf{x}^{(1)}\mathbf{x}^{(2)}$$

SVM Example: XOR Problem

$$g(\mathbf{x}) = -2x^{(1)}x^{(2)}$$



decision boundaries nonlinear



decision boundary is linear

SVM Summary

- Advantages:
 - Based on very strong theory
 - Excellent generalization properties
 - Objective function has no local minima
 - Can be used to find non linear discriminant functions
 - Complexity of the classifier is characterized by the number of support vectors rather than the dimensionality of the transformed space
- Disadvantages:
 - Directly applicable to two-class problems
 - Quadratic programming is computationally expensive
 - Need to choose kernel

Multi-Class SVMs

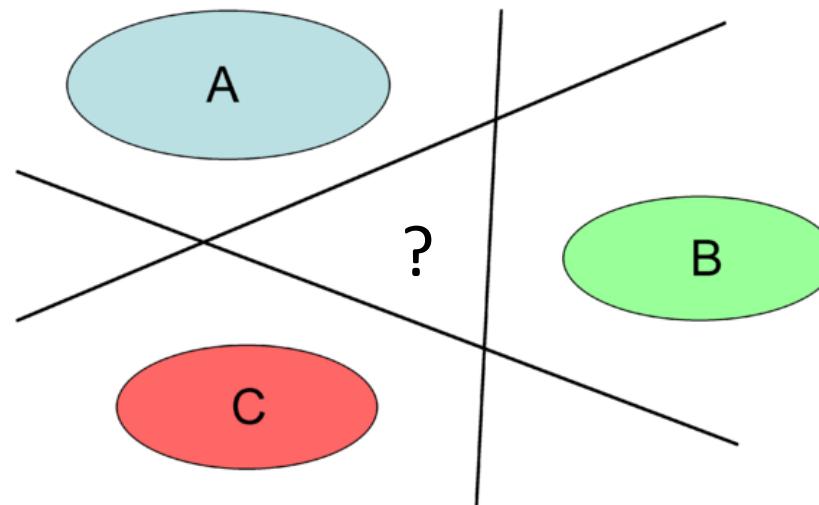
- One against all
- Pairwise
- These ideas apply to all binary classifiers when faced with multi-class problems

One-Against-All

- SVMs can only handle two-class outputs
- What can be done?
- Answer: learn N SVM's
 - SVM 1 learns "Output==1" vs "Output != 1"
 - SVM 2 learns "Output==2" vs "Output != 2"
 - ...
 - SVM N learns "Output==N" vs "Output != N"

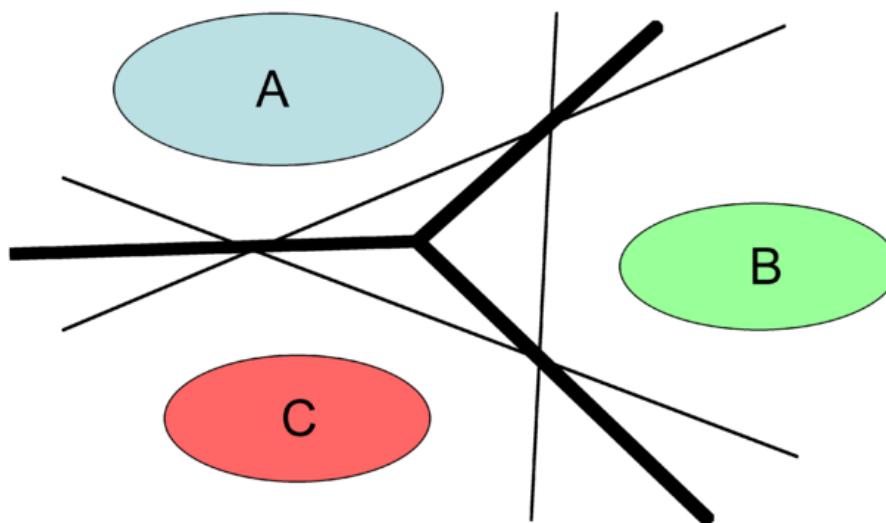
One-Against-All

- Original idea (Vapnik, 1995): classify x as ω_i if and only if the corresponding SVM accepts x and all other SVMs reject it



One-Against-All

- Modified idea (Vapnik, 1998): classify x according to the SVM that produces the highest value (use more than sign of decision function)



Pairwise SVMs

- Learn $N(N-1)/2$ SVM's
 - SVM 1 learns "Output==1" vs "Output == 2"
 - SVM 2 learns "Output==1" vs "Output == 3"
 - ...
 - SVM M learns "Output==N-1" vs "Output == N"

Pairwise SVMs

- To classify a new input, apply each SVM and choose the label that “wins” most often

